

# Introduction to Modern Fortran

*See next foil for copyright information*

Nick Maclaren

**nmm1@cam.ac.uk**

March 2014

# Acknowledgement

Derived from the course written and owned by

Dr. Steve Morgan  
Computing Services Department  
The University of Liverpool

The Copyright is joint between the authors  
Permission is required before copying it

Please ask if you want to do that

# Important!

There is a **lot** of material in the course  
And there is even more in extra slides ...

**Some** people will already know **some** Fortran  
**Some** will be programmers in **other** languages  
**Some** people will be complete **newcomers**

The course is intended for **all** of those people

- **Please** tell me if I am going too fast  
Not afterwards, but **as soon as** you have trouble

# Beyond the Course (1)

<http://people.ds.cam.ac.uk/nmm1/...>  
.../Fortran/  
.../OldFortran/  
.../Arithmetic/ etc.

## Programming in Fortran 90/95

by **Steve Morgan** and **Lawrie Schonfelder**  
(Fortran Market, PDF, **\$15**)

<http://www.fortran.com/>

Also **Fortran 90** version of that

# Beyond the Course (2)

## Fortran 95/2003 Explained

by Michael Metcalf, John Reid and  
Malcolm Cohen

Also Fortran 90 version of that

## Fortran 90 Programming

by Miles Ellis, Ivor Phillips and  
Thomas Lahey

# Beyond the Course (3)

SC22WG5 (ISO Fortran standard committee)

<http://www.nag.co.uk/sc22wg5/>

<http://www.fortran.com/fortran/>

⇒ ‘Information’, ‘Standards Documents’

Miscellaneous information and useful guidance

<http://www.star.le.ac.uk/~cgp/fortran.html>

Liverpool Course

<http://www.liv.ac.uk/HPC/...>

[.../HTMLFrontPageF90.html](http://www.liv.ac.uk/HPC/.../HTMLFrontPageF90.html)

## Beyond the Course (4)

A real, live (well coded) Fortran 95 application

<http://www.wannier.org>

Most of the others I have seen are not public  
Please tell me of any you find that are

# Important!

There is a **lot** of material in the course  
And there is even more in extra slides ...

This has been stripped down to the **bare minimum**  
Some **loose ends** will remain, unfortunately  
You will need to skip a **few** of the practicals

- **Please** tell me if I am going too fast  
Not afterwards, but **as soon as** you have trouble



# Practicals

These will be delayed until after second lecture  
Then there will be two practicals to do

One is using the compiler and diagnostics  
Just to see what happens in various cases

The other is questions about the basic rules

Full instructions will be given then  
Including your identifiers and passwords

# History

FORmula TRANslation invented 1954–8  
by John Backus and his team at IBM

FORTRAN 66 (ISO Standard 1972)

FORTRAN 77 (1980)

Fortran 90 (1991)

Fortran 95 (1996)

Fortran 2003 (2004)

Fortran 2008 (2011)

The “Old Fortran” slides have more detail

# Hardware and Software

A system is built from **hardware** and **software**

The **hardware** is the physical medium, e.g.

- CPU, memory, keyboard, display
- disks, ethernet interfaces etc.

The **software** is a set of computer programs, e.g.

- operating system, compilers, editors
- Fortran 90 programs

# Programs

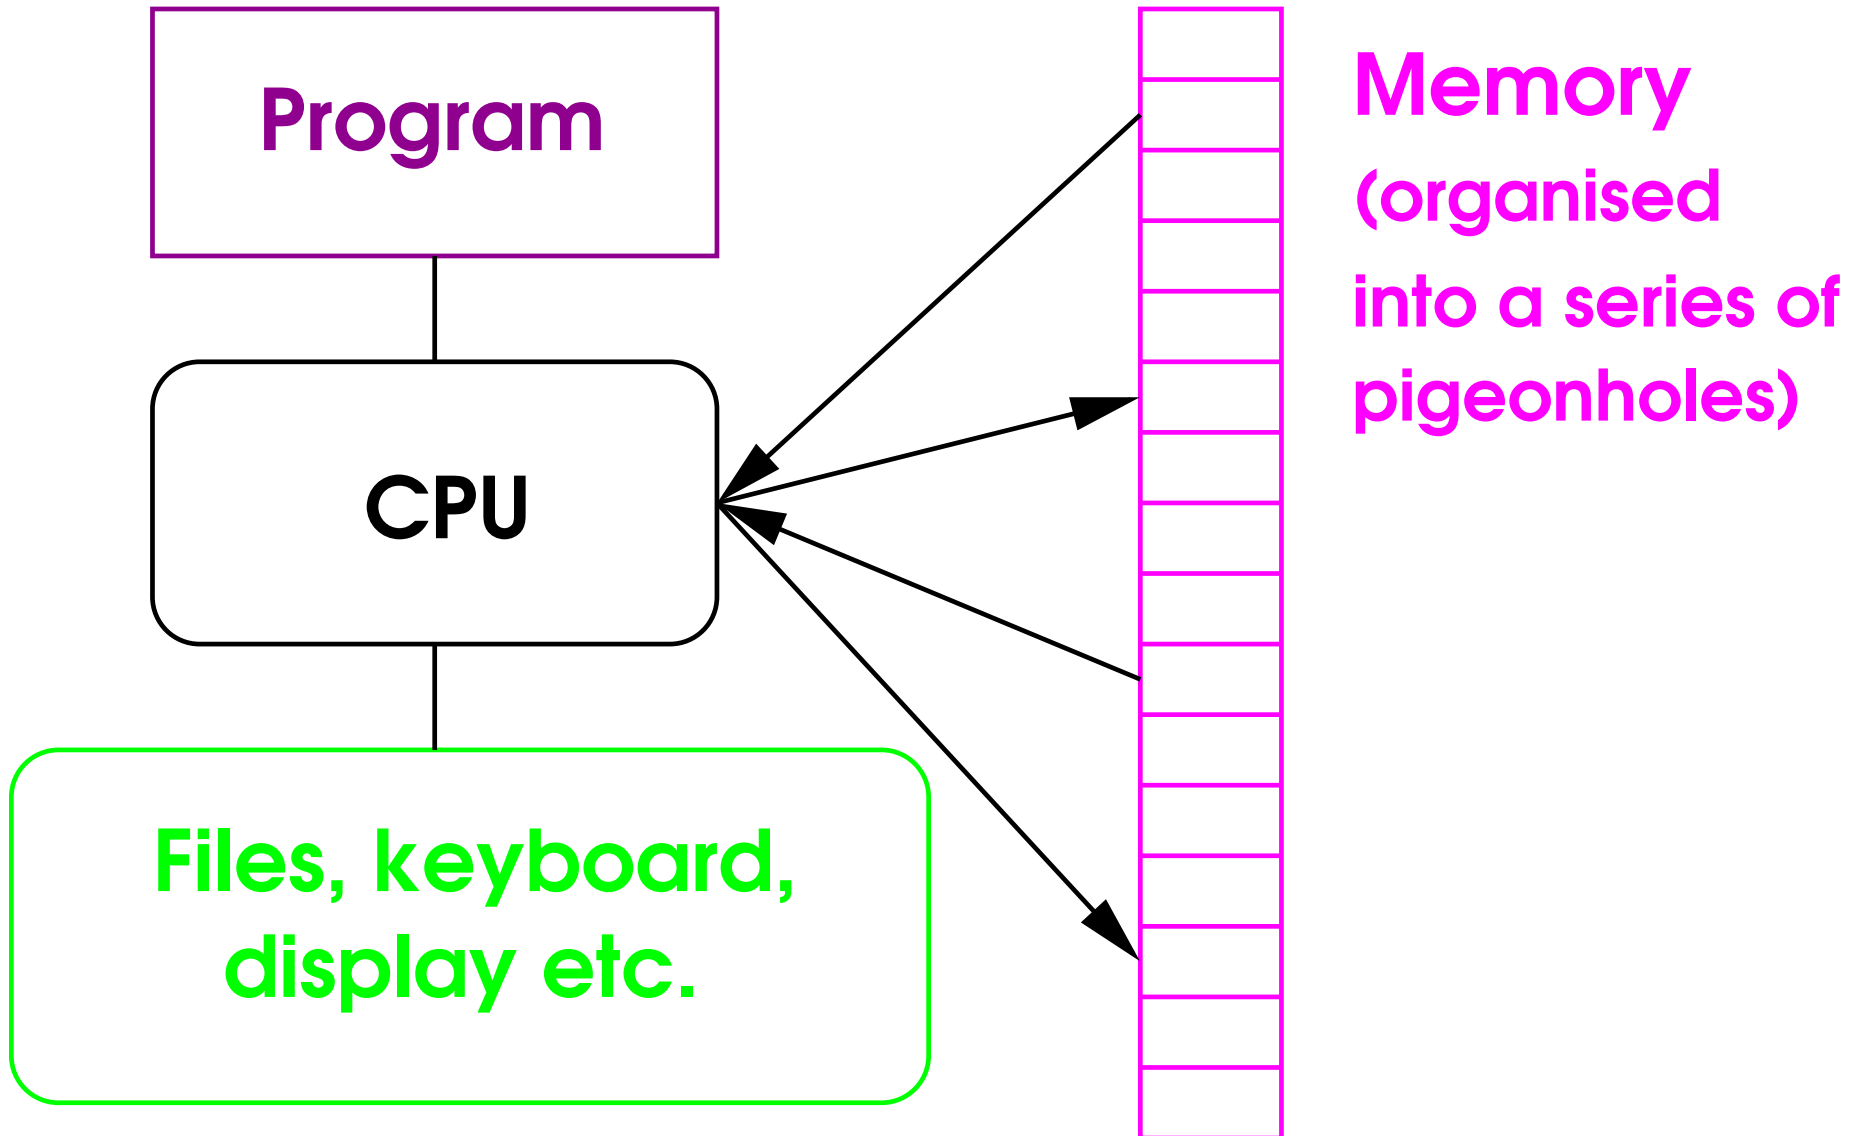
Fortran 90 is a **high-level language**  
Sometimes called “**third-generation**” or **3GL**

Uses English-like words and math-like expressions

```
Y = X + 3  
PRINT *, Y
```

**Compilers** translate into machine instructions  
A **linker** then creates an **executable program**  
The **operating system** runs the **executable**

# Fortran Programming Model



# Algorithms and Models

An **algorithm** is a set of **instructions**  
They are executed in **a defined order**  
Doing that carries out a specific task

The above is **procedural programming**  
Fortran 90 is a **procedural language**

**Object-orientation** is still procedural  
Fortran 90 has **object-oriented** facilities

# An Example of a Problem

Write a program to convert a time in hours, minutes and seconds to one in seconds

Algorithm:

1. Multiply the hours by 60
2. Add the minutes to the result
3. Multiply the result by 60
4. Add the seconds to the result

# Logical Structure

1. Start of program
2. Reserve memory for data
3. Write prompt to display
4. Read the time in hours, minutes and seconds
5. Convert the time into seconds
6. Write out the number of seconds
7. End of program



# The Program

```
PROGRAM example1
```

```
! Comments start with an exclamation mark
```

```
    IMPLICIT NONE
```

```
    INTEGER :: hours, mins, secs, temp
```

```
    PRINT *, 'Type the hours, minutes and seconds'
```

```
    READ *, hours, mins, secs
```

```
    temp = 60* ( hours*60 + mins ) + secs
```

```
    PRINT *, 'Time in seconds =', temp
```

```
END PROGRAM example1
```

# High Level Structure

1. Start of program (or procedure)  
`PROGRAM example1`
2. Followed by the specification part  
`declare types and sizes of data`
- 3–6. Followed by the execution part  
`all of the 'action' statements`
7. End of program (or procedure)  
`END PROGRAM example1`

Comments do nothing and can occur anywhere  
`! Comments start with an exclamation mark`

# Program and File Names

- The **program** and **file** names are not related  
**PROGRAM QES** can be in file **QuadSolver.f90**  
Similarly for most other Fortran components

**Some implementations** like the same names  
Sometimes converted to lower- or upper-case

The **compiler** documentation **should** tell you  
It is sometimes in the **system** documentation  
Please ask for help, but it is outside this course

# The Specification Part

## 2. Reserve memory for data

`INTEGER :: hours, mins, secs, temp`

`INTEGER` is the `type` of the variables

`hours`, `mins`, `secs` are used to hold input

The values read in are called the `input data`

`temp` is called a `workspace variable`

also called a `temporary variable` etc.

The `output data` are '`Time . . . =`' and `temp`

They can be any expression, not just a variable

# The Execution Part

3. Write prompt to display

```
PRINT *, 'Type the hours, ...'
```

4. Read the time in hours, minutes and seconds

```
READ *, hours, mins, secs
```

5. Convert the time into seconds

```
temp = 60*( hours*60 + mins) + secs
```

6. Write out the number of seconds

```
PRINT *, 'Time in seconds =', temp
```

# Assignment and Expressions

```
temp = 60*( hours*60 + mins) + secs
```

The RHS is a pseudo-mathematical **expression**  
It calculates the value to be stored

- Expressions are very like A-level formulae  
**Fortran** is **FOR**mula **TRAN**slation – remember?  
We will come to the detailed rules later

- **temp** = stores the value in the **variable**  
A **variable** is a memory cell in Fortran's model

# Really Basic I/O

READ \*, <variable list> reads from **stdin**

PRINT \*, <expression list> writes to **stdout**

Both do input/output as **human-readable text**

Each I/O **statement** reads/writes on a new line

A **list** is **items** separated by **commas** (‘,’)

**Variables** are anything that can store **values**

**Expressions** are anything that deliver a **value**

Everything else will be explained later

# Repeated Instructions

The previous program handled only one value  
A more flexible one would be:

1. Start of program
2. Reserve memory for data
3. Repeat this until end of file
  - 3.1 Read the value of seconds
  - 3.2 Convert to minutes and seconds
  - 3.3 Write out the result
4. End of Program



# Sequences and Conditionals

Simple algorithms are just sequences

A simple algorithm for charging could be:

1. Calculate the bill
2. Print the invoice

Whereas it probably should have been:

1. Calculate the bill
2. If the bill exceeds minimum
  - 2.1 Then print the invoice
3. Otherwise
  - 3.1 Add bill to customer's account

# Summary

There are three basic control structures:

- A simple sequence
- A conditional choice of sequences
- A repeated sequence

All algorithms can be expressed using these

In practice, other structures are convenient

Almost always need to split into simpler tasks

Even Fortran II had subroutines and functions!

Doing that is an important language-independent skill

# Developing a Computer Program

There are four main steps:

1. Specify the problem
2. Analyse and subdivide into tasks
3. Write the Fortran 90 code
4. Compile and run (i.e. test)

Each step may require several iterations

You may need to restart from an earlier step

- The testing phase is **very** important

# Errors

- If the **syntax** is incorrect, the compiler says so  
For example: **INTEGER :: ,mins, secs**

- If the action is **invalid**, things are messier  
For example: **X/Y** when **Y** is zero  
**/** represents **division**, because of the lack of **÷**

You may get an error message at run-time

The program may crash, just stop or hang

It may produce nonsense values or go haywire

# Introduction to Modern Fortran

## *Fortran Language Rules*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Coverage

This course is modern, **free-format** source only  
[ If you don't understand this, don't worry ]  
The same applies to features covered later

Almost all **old Fortran** remains legal  
Avoid using it, as modern Fortran is better  
This mentions old Fortran only in passing

See the **OldFortran** course for those aspects  
It describes **fixed-format** and conversion  
Or ask questions or for help on such things, too

# Important Warning

Fortran's **syntax** is verbose and horrible  
It can fairly be described as a historical mess  
Its **semantics** are fairly clean and consistent

Its verbosity causes problems for examples  
Many of them use poor style, to be readable  
And they mostly omit essential error checking

- Do what I say, don't do what I do

Sorry about that . . .

# Correctness

Humans understand language quite well even when it isn't strictly correct

Computers (i.e. compilers) are not so forgiving

- **Programs** must follow the rules to the letter
- Fortran compilers **will** flag **all syntax** errors

Good compilers will detect more than is required

But **your** error may just change the meaning

Or do something invalid (“**undefined behaviour**”)



# Examples of Errors

Consider  $(N * M / 1024 + 5)$

If you mistype the '0' as a ')':  $(N * M / 1)24 + 5)$

You will get an error message when compiling  
It may be confusing, but will point out a problem

If you mistype the '0' as a '-':  $(N * M / 1 - 24 + 5)$

You will simply evaluate a different formula  
And get wrong answers with no error message

And if you mistype '\*' as '8'?

# Character Set

Letters (A to Z and a to z) and digits (0 to 9)  
Letters are matched ignoring their case

And the following special characters

\_ = + - \* / ( ) , . ' : ! " % & ; < > ? \$

Plus space (i.e. a blank), but not tab

The end-of-line indicator is not a character

Any character allowed in comments and strings

- Case is significant in strings, and only there

# Special Characters

`_ = + - * / ( ) , . ' : ! " % & ; < > ? $`

slash (/) is also used for divide

hyphen (−) is also used for minus

asterisk (\*) is also used for multiply

apostrophe (') is used for single quote

period (.) is also used for decimal point

The others are described when we use them

# Layout

- Do **not** use **tab**, **form-feed** etc. in your source  
Use no positioning except **space** and line breaks

Compilers do bizarre things with anything else  
Will work with some compilers but not others  
And can produce some very strange output

Even in **C**, using them is a recipe for confusion  
The really masochistic should ask me offline

# Source Form (1)

Spaces are not allowed in **keywords** or **names**  
**INTEGER** is not the same as **INT** **EGER**

**HOURS** is the same as **hoURs** or **hours**  
But not **HO** **URS** – that means **HO** and **URS**

- Some **keywords** can have two forms  
E.g. **ENDDO** is the same as **END** **DO**  
But **EN** **DDO** is treated as **EN** and **DDO**

⇒ **END** **DO** etc. is the direction Fortran is going

## Source Form (2)

- Do not run **keywords** and **names** together

**INTEGER****I,J,K** – **illegal**

**INTEGER** **I,J,K** – **allowed**

- You can use spaces liberally for clarity

**INTEGER** **I** , **J** , **K**

Exactly **where** you use them is a matter of taste

- Blank lines can be used in the same way  
Or lines consisting only of comments

# Double Colons

For descriptive names use **underscore**  
**largest\_of**, **maximum\_value** or **P12\_56**

- Best to use a double colon in **declarations**  
Separates **type specification** from **names**  
**INTEGER :: I, J, K**

This form is essential where **attributes** are used  
**INTEGER, INTENT(IN) :: I, J, K**

# Lines and Comments

A **line** is a sequence of up to **132** characters

A **comment** is from **!** to the end of line

The whole of a comment is totally ignored

**A = A+1    ! These characters are ignored**

**! That applies to !, & and ; too**

Blank lines are completely ignored

**!**

**! Including ones that are just comments**

**!**



# Use of Layout

Well laid-out programs are much more readable  
You are less likely to make trivial mistakes  
And **much** more likely to spot them!

This also applies to **low-level** formats, too  
E.g. **1.0e6** is clearer than **1.e6** or **.1e7**

- None of this is Fortran-specific

# Use of Comments

Appropriate commenting is very important

This course does **not** cover that topic

And, often, comments are omitted for brevity

“How to Help Programs Debug Themselves”

Gives guidelines on how best to use comments

- This isn't Fortran-specific, either

# Use of Case

- Now, this **IS** Fortran-specific!

It doesn't matter what case convention you use

- But **DO** be **moderately**† consistent!

Very important for clarity and editing/searching

For example:

**UPPER** case for keywords, **lower** for names

You may prefer **Capitalised** names

† *A foolish consistency is the hobgoblin of little minds*

# Statements and Continuation

- A **program** is a sequence of **statements**

Used to build high-level constructs

**Statements** are made up out of **lines**

- Statements are continued by appending &

```
A = B + C + D + E + &  
      F + G + H
```

Is equivalent to

```
A = B + C + D + E + F + G + H
```

# Other Rules (1)

**Statements** can start at any position

- Use indentation to clarify your code

```
IF (a > 1.0) THEN
    b = 3.0
ELSE
    b = 2.0
END IF
```

- A number starting a statement is a **label**

```
10 A = B + C
```

The use of **labels** is described later

## Other Rules (2)

You can put **multiple statements** on a line

```
a = 3 ; b = 4 ; c = 5
```

Overusing that can make a program unreadable

But it can clarify your code in some cases

Avoid mixing **continuation** with that or **comments**

It works, but can make code very hard to read

```
a = b + c ; d = e + f + &  
g + h
```

```
a = b + c + & ! More coming ...  
d = e + f + g + h
```

# Breaking Character Strings

- Continuation lines can start with an &  
Preceding spaces and the & are suppressed

The following works and allows indentation:

```
PRINT 'Assume that this string &  
      &is far too long and complic&  
      &ated to fit on a single line'
```

The initial & avoids including excess spaces  
And avoids problems if the text starts with !

This may also be used to continue any line

# Names

Up to 31 (now 63) letters, digits and underscores

- Names must start with a letter

Upper and lower case are equivalent

DEPTH, Depth and depth are the same name

The following are valid Fortran names

A, AA, aaa, Tax, INCOME, Num1, NUM2, NUM333,  
N12MO5, atmospheric\_pressure, Line\_Colour,  
R2D2, A\_21\_173\_5a



# Invalid Names

The following are **invalid names**

**1A** does not begin with a **letter**

**\_B** does not begin with a **letter**

**Depth\$0** contains an illegal character '\$'

**A-3** would be interpreted as subtract **3** from **A**

**B.5:** illegal characters '.' and ':'

**A\_name\_made\_up\_of\_more\_than\_31\_letters**  
too long, **38** characters

# Compiling and Testing

We shall use the **gfortran** under **Linux**  
PWF/MCS/DS **Windows** does not have a Fortran  
Using any Fortran compiler is much the same

Please ask about anything you don't understand  
Feel free to bring problems with other Fortrans  
Feel free to use **gdb** if you know it

Solutions to exercises available from  
<http://people.ds.cam.ac.uk/nmm1/Fortran/Answers>

# Instructions

If running **Microsoft Windows**, **CTRL-ALT-DEL**

Select **Restart** and then **Linux**

Log into **Linux** and start a shell and an editor

Create programs called **prog.f90**, **fred.f90** etc.

- Run by typing commands like  
**nagfor -C=all -o fred fred.f90**  
**./fred**
- Analyse what went wrong
- Fix bugs and retry

# Instructions

- Run by typing commands like  
gfortran -g -O3 -Wall -Wextra -o fred fred.f90  
./fred
- Analyse what went wrong
- Fix bugs and retry

# Introduction to Modern Fortran

## *Data Types and Basic Calculation*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Data Types (1)

- **INTEGER** for exact whole numbers

e.g. 1, 100, 534, -18, -654321 etc.

In maths, an approximation to the ring  $\mathbb{Z}$

- **REAL** for approximate, fractional numbers

e.g. 1.1, 3.0, 23.565,  $\pi$ ,  $\exp(1)$ , etc.

In maths, an approximation to the field  $\mathbb{R}$

- **COMPLEX** for complex, fractional numbers

e.g. (1.1, -23.565), etc.

In maths, an approximation to the field  $\mathbb{C}$

## Data Types (2)

- **LOGICAL** for truth values

These may have only values **true** or **false**

e.g. **.TRUE.** , **.FALSE.**

These are often called **boolean** values

- **CHARACTER** for strings of characters

e.g. **'?'**, **'Albert Einstein'**, **'X + Y = '**, etc.

The **string length** is part of the **type** in Fortran

There is no separate **character type**, unlike **C**

There is more on this later

# Integers (1)

- **Integers** are restricted to lie in a finite range

Typically  $\pm 2147483647$  ( $-2^{31}$  to  $2^{31}-1$ )

Sometimes  $\pm 9.23 \times 10^{17}$  ( $-2^{63}$  to  $2^{63}-1$ )

A compiler **may** allow you to select the range

Often including  $\pm 32768$  ( $-2^{15}$  to  $2^{15}-1$ )

Older and future systems may have other ranges

There is more on the arithmetic and errors later



# Integers (2)

Fortran uses **integers** for:

- **Loop counts** and **loop limits**
- An **index** into an **array** or a position in a list
- An **index** of a **character** in a **string**
- As **error codes**, **type categories** etc.

Also use them for purely **integral values**

E.g. calculations involving **counts** (or money)

They can even be used for **bit masks** (see later)

# Integer Constants

Usually, an optional **sign** and one or more **digits**

e.g. 0, 123, -45, +67, 00012345

E.g. '60' in `minutes = minutes + 60*hours`

Also allowed in **binary**, **octal** and **hexadecimal**

e.g. `B'011001'`, `O'35201'`, `Z'a12bd'`

- As with names, the case is irrelevant

There is a little more, which is covered later

# Reals

- **Reals** are held as **floating-point** values  
These also have a finite **range** **and** **precision**

It is essential to use **floating-point** appropriately

- Much of the Web is misleading about this

This course will mention only the bare minimum

See “**How Computers Handle Numbers**”

There is simplified version of that later on

**Reals** are used for **continuously varying** values

Essentially just as you were taught at A-level

# IEEE 754

You can assume a variant of IEEE 754

You should almost always use IEEE 754 64-bit

There is information on how to select it later

IEEE 754 32-, 64- and 128-bit formats are:

$10^{-38}$  to  $10^{+38}$  and 6–7 decimal places

$10^{-308}$  to  $10^{+308}$  and 15–16 decimal places

$10^{-4932}$  to  $10^{+4932}$  and 33–34 decimal places

Older and future systems may be different

# Real Constants

- Real constants **must** contain a **decimal point** or an **exponent**

They can have an optional **sign**, just like **integers**

The basic **fixed-point form** is anything like:

123.456, -123.0, +0.0123, 123., .0123  
0012.3, 0.0, 000., .000

- Optionally followed **E** or **D** and an exponent  
1.0E6, 123.0D-3, .0123e+5, 123.d+06, .0e0

1E6 and 1D6 are also valid Fortran **real constants**

# Complex Numbers

This course will generally ignore them  
If you don't know what they are, don't worry

These are (real, imaginary) pairs of REALs  
I.e. Cartesian notation, as at A-level

Constants are pairs of reals in parentheses  
E.g. (1.23, -4.56) or (-1.0e-3, 0.987)

# Declaring Numeric Variables

Variables hold values of different types

`INTEGER :: count, income, mark`

`REAL :: width, depth, height`

You can get all undeclared variables diagnosed  
Add the statement `IMPLICIT NONE` at the start  
of every `program`, `subroutine`, `function` etc.

If not, variables are declared implicitly by use

Names starting with `I–N` are `INTEGER`

Ones with `A–H` and `O–Z` are `REAL`

# Implicit Declaration

- This is a common source of errors

`REAL :: metres, inches`

`inch3s = meters*30.48`

The effects can be even worse for **function calls**

Ask offline if you want to know the details

- Also the default `REAL` type is a disaster

Too inaccurate for practical use (see later)

- You should **always** use `IMPLICIT NONE`



# Important Warning!

- I shall **NOT** do that myself

I warned you about this in the previous lecture  
The problem is fitting all the text onto a slide  
I shall often rely on **implicit typing** :-(

- **Do what I say, don't do what I do**

If I omit it in example files, it is a **BUG**

# Assignment Statements

The general form is

**<variable> = <expression>**

This is actually very powerful (see later)

This **first** evaluates the **expression** on the **RHS**

It **then** stores the result in the **variable** on the **LHS**

It replaces whatever **value** was there before

For example:

**Max = 2 \* Min**

**Sum = Sum + Term1 + Term2 + (Eps \* Err)**

# Arithmetic Operators

There are five built-in numeric operations

- + addition
  - − subtraction
  - \* multiplication
  - / division
  - \*\* exponentiation (i.e. raise to the power of)
- Exponents can be any arithmetic type:  
INTEGER, REAL or COMPLEX
- Generally, it is best to use them in that order

# Examples

Some examples of arithmetic expressions are

$A * B - C$

$A + C1 - D2$

$X + Y / 7.0$

$2 ** K$

$A ** B + C$

$A * B - C$

$(A + C1) - D2$

$A + (C1 - D2)$

$P ** 3 / ((X + Y * Z) / 7.0 - 52.0)$

# Operator Precedence

Fortran uses normal mathematical conventions

- Operators bind according to **precedence**
- And then generally, from **left to right**

The **precedence** from highest to lowest is

**\*\***    **exponentiation**

**\*** /    **multiplication and division**

**+** **-**    **addition and subtraction**

- Parentheses (‘**(**’ and ‘**)**’) are used to control it  
Use them whenever the **order matters** or it is **clearer**

# Examples

$X + Y * Z$  is equivalent to  $X + (Y * Z)$

$X + Y / 7.0$  is equivalent to  $X + (Y / 7.0)$

$A - B + C$  is equivalent to  $(A - B) + C$

$A + B ** C$  is equivalent to  $A + (B ** C)$

$- A ** 2$  is equivalent to  $-(A ** 2)$

$A - (((B + C)))$  is equivalent to  $A - (B + C)$

- You can force any order you like

$(X + Y) * Z$

Adds  $X$  to  $Y$  and **then** multiplies by  $Z$

# Warning

$X + Y + Z$  may be evaluated as any of  
 $X + (Y + Z)$  or  $(X + Y) + Z$  or  $Y + (X + Z)$  or . . .

Fortran defines **what** an expression means  
It does not define **how** it is calculated

They are all **mathematically** equivalent  
But may sometimes give slightly different results

See “**How Computers Handle Numbers**” for more

# Precedence Problems

Mathematical conventions vary in some aspects

$A / B * C$  – is it  $A / (B * C)$  or  $(A / B) * C$ ?

$A ** B ** C$  – is it  $A ** (B ** C)$  or  $(A ** B) ** C$ ?

Fortran specifies that:

$A / B * C$  is equivalent to  $(A / B) * C$

$A ** B ** C$  is equivalent to  $A ** (B ** C)$

- Yes, **\*\*** binds from right to left!



# Parenthesis Problems

**Always** use **parentheses** in ambiguous cases  
If only to imply “**Yes, I really meant that**”

And to help readers used to different rules  
Programming languages vary in what they do

Be careful of over-doing it – what does this do?

$$(((A+(P*R+B)/2+B**3)/(4/Y)*C+D))+E)$$

- Several, simpler statements is better style

# Integer Expressions

I.e. ones of **integer constants** and **variables**

```
INTEGER :: K, L, M  
N = K*(L+2)/M**3-N
```

These are evaluated in **integer arithmetic**

- Division always truncates **towards zero**

If  $K = 4$  and  $L = 5$ , then  $K+L/2$  is 6  
 $(-7)/3$  and  $7/(-3)$  are both -2

# Mixed Expressions

**INTEGER** and **REAL** is evaluated as **REAL**  
Either and **COMPLEX** goes to **COMPLEX**

Be careful with this, as it can be deceptive

```
INTEGER :: K = 5  
REAL :: X = 1.3  
X = X + K/2
```

That will add **2.0** to **X**, not **2.5**  
**K/2** is still an **INTEGER** expression

# Conversions

There are several ways to force conversion

- Intrinsic functions **INT**, **REAL** and **COMPLEX**

$X = X + \text{REAL}(K)/2$

$N = 100 * \text{INT}(X/1.25) + 25$

You can use appropriate constants

You can even add zero or multiply by one

$X = X + K/2.0$

$X = X + (K + 0.0)/2$

The last isn't very nice, but works well enough

And see later about **KIND** and precision

# Mixed-type Assignment

<real variable> = <integer expression>

- The RHS is converted to **REAL**  
Just as in a mixed-type expression

<integer variable> = <real expression>

- The RHS is **truncated** to **INTEGER**  
It is always truncated **towards zero**

Similar remarks apply to **COMPLEX**

- The **imaginary** part is discarded, quietly

The **RHS** is evaluated **independently** of the **LHS**

# Examples

```
INTEGER :: K = 9, L = 5, M = 3, N
```

```
REAL :: X, Y, Z
```

```
X = K ; Y = L ; Z = M
```

```
N = (K/L)*M
```

```
N = (X/Y)*Z
```

**N** will be 3 and 5 in the two cases

$(-7)/3 = 7/(-3) = -2$  and  $7/3 = (-7)/(-3) = 2$

# Numeric Errors

See “How Computers Handle Numbers”

This is a **very** minimal summary

- **Overflowing the range** is a serious error  
As is **dividing by zero** (e.g. **123/0** or **0.0/0.0**)  
Fortran does not define what those cases do
- Each **numeric type** may behave differently  
Even different compiler options will, too
- And do **not** assume results are predictable

# Examples

Assume the **INTEGER** range is  $\pm 2147483647$

And the **REAL** range is  $\pm 10^{38}$

- Do you know what this is defined to do?

```
INTEGER :: K = 1000000  
REAL :: X = 1.0e20  
PRINT *, (K*K)/K, (X*X)/X
```

- The answer is “**anything**” – and it means it  
Compilers optimise on the basis of no errors  
**Numeric errors** can cause **logical errors**



# Numeric Non-Errors (1)

- Conversion to a **lesser** type loses information  
You will get no warning of this, unfortunately

**REAL**  $\Rightarrow$  **INTEGER** truncates towards zero

**COMPLEX**  $\Rightarrow$  **REAL** drops the imaginary part

**COMPLEX**  $\Rightarrow$  **INTEGER** does both of them

That also applies when dropping in **precision**

E.g. assigning a **64-bit** real to a **32-bit** one

# Numeric Non-Errors (2)

Fortran does **NOT** specify the following  
But it is true on all systems you will use

Results too small to represent are not errors

- They will be replaced by zero if necessary
- **Inexact results** round to the nearest value

That also applies when dropping in **precision**

# Intrinsic Functions

Built-in functions that are always available

- No declaration is needed – just use them

For example:

```
Y = SQRT(X)
```

```
PI = 4.0 * ATAN(1.0)
```

```
Z = EXP(3.0*Y)
```

```
X = REAL(N)
```

```
N = INT(X)
```

```
Y = SQRT(-2.0*LOG(X))
```

# Intrinsic Numeric Functions

REAL(n)      ! Converts its argument n to REAL  
INT(x)        ! Truncates x to INTEGER (to zero)  
AINT(x)       ! The result remains REAL  
NINT(x)       ! Converts x to the nearest INTEGER  
ANINT(x)      ! The result remains REAL  
ABS(x)        ! The absolute value of its argument  
! Can be used for INTEGER, REAL or COMPLEX  
MAX(x,y,...) ! The maximum of its arguments  
MIN(x,y,...) ! The minimum of its arguments  
MOD(x,y)      ! Returns x modulo y

And there are more – some are mentioned later

# Intrinsic Mathematical Functions

SQRT(x)    ! The square root of x  
EXP(x)     ! e raised to the power x  
LOG(x)     ! The natural logarithm of x  
LOG10(x)   ! The base 10 logarithm of x

SIN(x)     ! The sine of x, where x is in radians  
COS(x)     ! The cosine of x, where x is in radians  
TAN(x)     ! The tangent of x, where x is in radians  
ASIN(x)    ! The arc sine of x in radians  
ACOS(x)    ! The arc cosine of x in radians  
ATAN(x)    ! The arc tangent of x in radians

And there are more – see the references

# Bit Masks

As in **C** etc., **integers** are used for these  
Use is by weirdly-named functions (historical)

Bit indices start at **zero**

Bit **K** has value  $2^K$  (little-endian)

As usual, stick to **non-negative** integers

- A little tedious, but very easy to use

# Bit Ininsics

BIT_SIZE(x)	! The number of bits in x
BTEST(x, n)	! Test bit n of x
IBSET(x, n)	! Set bit n of x
IBCLR(x, n)	! Clear bit n of x
IBITS(x, m, n)	! Extract n bits
NOT(x)	! NOT x
IAND(x, y)	! x AND y
IOR(x, y)	! x OR y
IEOR(x, y)	! x (exclusive or) y
ISHFT(x, n)	! Logical shift
ISHFTC(x, n, [k])	! Circular shift

# Logical Type

These can take only two values: **true** or **false**  
**.TRUE.** and **.FALSE.**

- Their type is **LOGICAL** (not **BOOL**)

**LOGICAL :: red, amber, green**

**IF (red) THEN**

**PRINT \*, 'Stop'**

**red = .False. ; amber = .True. ; green = .False.**

**ELSIF (red .AND. amber) THEN**

**. . .**



# Relational Operators

- Relations create LOGICAL values

These can be used on any other built-in type

== (or .EQ.) equal to

/= (or .NE.) not equal to

These can be used only on INTEGER and REAL

< (or .LT.) less than

<= (or .LE.) less than or equal

> (or .GT.) greater than

>= (or .GE.) greater than or equal

See “How Computers Handle Numbers” for more

# Logical Expressions

Can be as complicated as you like

Start with `.TRUE.`, `.FALSE.` and `relations`

Can use `parentheses` as for numeric ones

`.NOT.`, `.AND.` and `.OR.`

`.EQV.` must be used instead of `==`

`.NEQV.` must be used instead of `/=`

- Fortran is `not` like `C`-derived languages  
`LOGICAL` is not a sort of `INTEGER`

# Short Circuiting

```
LOGICAL :: flag
```

```
flag = ( Fred( ) > 1.23 .AND. Joe( ) > 4.56 )
```

Fred and Joe may be called in either order

If Fred returns 1.1, then Joe may not be called

If Joe returns 3.9, then Fred may not be called

Fortran expressions define the answer only

The behaviour is up to the compiler

One of the reasons that it is so optimisable

# Character Type

Used when **strings of characters** are required  
Names, descriptions, headings, etc.

- Fortran's basic type is a **fixed-length string**  
Unlike almost all more recent languages

- **Character constants** are **quoted strings**

**PRINT \*, 'This is a title'**

**PRINT \*, "And so is this"**

The **characters** between **quotes** are the **value**

# Character Data

- The case of letters is significant in them  
Multiple spaces are not equivalent to one space  
Any representable character may be used

The only Fortran syntax where the above is so  
Remember the line joining method?

In 'Time^^=^^13:15', with '^' being a space  
The character string is of length 15  
Character 1 is T, 8 is a space, 10 is 1 etc.

# Character Constants

- "This has UPPER, lower and MiXed cases"
- 'This has a double quote (") character'
- "Apostrophe (') is used for single quote"
- "Quotes (\"") are escaped by doubling"
- 'Apostrophes (') are escaped by doubling'
- 'ASCII ' , | , ~ , ^ , @ , # and \ are allowed here'
- "Implementations may do non-standard things"
- 'Backslash (\) MAY need to be doubled'
- "Avoid newlines, tabs etc. for your own sanity"

# Character Variables

```
CHARACTER :: answer, marital_status  
CHARACTER(LEN=10) :: name, dept, faculty  
CHARACTER(LEN=32) :: address
```

`answer` and `marital_status` are each of length **1**

They hold precisely **one character** each

`answer` might be **blank**, or hold **'Y'** or **'N'**

`name`, `dept` and `faculty` are of length **10**

And `address` is of length **32**

# Another Form

```
CHARACTER :: answer*1,      &  
           marital_status*1, name*10,    &  
           dept*10, faculty*10, address*32
```

While this form is historical, it is more compact

- Don't mix the forms – this is an abomination

```
CHARACTER(LEN=10) :: dept, faculty, addr*32
```

- For obscure reasons, using **LEN=** is cleaner

It avoids some arcane syntactic “gotchas”



# Character Assignment

```
CHARACTER(LEN=6) :: forename, surname  
forename = 'Nick'  
surname = 'Maclaren'
```

`forename` is padded with spaces ('Nick^^')  
`surname` is truncated to fit ('Maclar')

- Unfortunately, you won't get told  
But at least it won't overwrite something else

# Character Concatenation

Values may be **joined** using the **//** operator

```
CHARACTER(LEN=6) :: identity, A, B, Z  
identity = 'TH' // 'OMAS'  
A = 'TH' ; B = 'OMAS'  
Z = A // B  
PRINT *, Z
```

Sets **identity** to **'THOMAS'**

But **Z** looks as if it is still **'TH'** – why?

**//** does not remove **trailing spaces**

It uses the whole length of its inputs

# Substrings

If `Name` has length `9` and holds `'Marmaduke'`

`Name(1:1)` would refer to `'M'`

`Name(2:4)` would refer to `'arm'`

`Name(6:)` would refer to `'duke'` – note the form!

We could therefore write statements such as

```
CHARACTER :: name*20, surname*18, title*4
```

```
name = 'Dame Edna Everage'
```

```
title = name(1:4)
```

```
surname = name(11:)
```

# Example

This is not an example of good style!

```
PROGRAM message
  IMPLICIT NONE
  CHARACTER :: mess*72, date*14, name*40
  mess = 'Program run on'
  mess(30:) = 'by'
  READ *, date, name
  mess(16:29) = date
  mess(33:) = name
  PRINT *, mess
END PROGRAM message
```

# Warning – a “Gotcha”

CHARACTER substrings look like array sections  
But there is no equivalent of array indexing

```
CHARACTER :: name*20, temp*1  
temp = name(10)
```

- name(10) is an implicit function call  
Use name(10:10) to get the tenth character

CHARACTER variables come in various lengths  
name is not made up of 20 variables of length 1

# Character Intrinsic

LEN(c)	! The STORAGE length of c
TRIM(c)	! c without trailing blanks
ADJUSTL(C)	! With leading blanks removed
INDEX(str,sub)	! Position of sub in str
SCAN(str,set)	! Position of any character in set
REPEAT(str,num)	! num copies of str, joined

And there are more – see the references

# Examples

```
name = '  Bloggs  '  
newname = TRIM(ADJUSTL(name))
```

**newname** would contain **'Bloggs'**

```
CHARACTER(LEN=6) :: A, B, Z  
A = 'TH' ; B = 'OMAS'  
Z = TRIM(A) // B
```

**Now Z** gets set to **'THOMAS'** correctly!

# Collation Sequence

This controls whether "ffred" < "Fred" or not

- Fortran is **not** a **locale**-based language  
It specifies **only** the following

'A' < 'B' < ... < 'Y' < 'Z'	These ranges
'a' < 'b' < ... < 'y' < 'z'	will not
'0' < '1' < ... < '8' < '9'	overlap
' ' is less than all of 'A', 'a' and '0'	

A **shorter** operand is extended with **blanks** (' ')



# Working with CHARACTER

This is one of the things that has been **omitted**  
It's there in the notes, if you are interested

Can **assign**, **concatenate** and **compare** them  
Can extract **substrings** and do lots more

But, for the academy, you don't need to do that

- Skip the practicals that need those facilities

# Named Constants (1)

- These have the **PARAMETER** attribute

```
REAL, PARAMETER :: pi = 3.14159  
INTEGER, PARAMETER :: maxlen = 100
```

They can be used anywhere a **constant** can be

```
CHARACTER(LEN=maxlen) :: string  
circum = pi * diam  
IF (nchars < maxlen) THEN  
    . . .
```

# Named Constants (2)

Why are these important?

They reduce mistyping errors in long numbers

Is 3.14159265358979323846D0 correct?

They can make formulae etc. much clearer

Much clearer **which** constant is being used

They make it easier to modify the program later

INTEGER, PARAMETER :: MAX\_DIMENSION = 10000

# Named Character Constants

```
CHARACTER(LEN=*), PARAMETER ::      &  
    author = 'Dickens', title = 'A Tale of Two Cities'
```

**LEN=\*** takes the length from the data

It is permitted to define the **length** of a constant  
The data will be **padded** or **truncated** if needed

- But the above form is generally the best

# Named Constants (3)

- Expressions are allowed in constant values

```
REAL, PARAMETER :: pi = 3.14135, &  
    pi_by_4 = pi/4, two_pi = 2*pi, &  
    e = exp(1.0)
```

```
CHARACTER(LEN=*), PARAMETER :: &  
    all_names = 'Pip, Squeak, Wilfred', &  
    squeak = all_names(6:11)
```

Generally, anything reasonable is allowed

- It must be determinable at compile time

# Initialisation

- Variables start with **undefined** values  
They often vary from run to run, too
- **Initialisation** is very like defining constants  
Without the **PARAMETER** attribute

```
INTEGER :: count = 0, I = 5, J = 100
REAL :: inc = 1.0E5, max = 10.0E5, min = -10.0E5
CHARACTER(LEN=10) :: light = 'Amber'
LOGICAL :: red = .TRUE., blue = .FALSE., &
           green = .FALSE.
```

# Information for Practicals

A **program** has the following basic **structure**:

**PROGRAM** name

Declarations

Other statements

**END PROGRAM** name

Read and write data from the terminal using:

**READ \***, variable [ , variable ]...

**PRINT \***, expression [ , expression ]...

# Introduction to Modern Fortran

## *Control Constructs*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014



# Control Constructs

These change the **sequential execution** order  
We cover the main **constructs** in some detail  
We shall cover **procedure call** later

The main ones are:

- Conditionals** (**IF** etc.)

- Loops** (**DO** etc.)

- Switches** (**SELECT/CASE** etc.)

- Branches** (**GOTO** etc.)

**Loops** are by far the most complicated

# Single Statement IF

Oldest and simplest is the single statement **IF**  
**IF (logical expression) simple statement**

If the **LHS** is **.True.**, the **RHS** is executed

If not, the **whole statement** has no effect

```
IF (MOD(count,1000) == 0) &  
    PRINT *, 'Reached', count
```

```
IF (X < A) X = A
```

Unsuitable for anything complicated

- Only **action statements** can be on the **RHS**

No **IFs** or statements containing **blocks**

# Block IF Statement

A **block IF** statement is more flexible

The following is the most traditional form of it

```
IF (logical expression) THEN
    then block of statements
ELSE
    else block of statements
END IF
```

If the **expr.** is **.True.**, the **first** block is executed

If not, the **second** one is executed

**END IF** can be spelled **ENDIF**

# Example

```
LOGICAL :: flip
```

```
IF (flip .AND. X /= 0.0) THEN  
    PRINT *, 'Using the inverted form'  
    X = 1.0/A  
    Y = EXP(-A)  
ELSE  
    X = A  
    Y = EXP(A)  
END IF
```

# Omitting the ELSE

The **ELSE** and its block can be omitted

```
IF (X > Maximum) THEN  
    X = Maximum  
END IF
```

```
IF (name(1:4) == "Miss" .OR.      &  
    name(1:4) == "Mrs.") THEN  
    name(1:3) = "Ms."  
    name(4:) = name(5:)  
END IF
```

# Including ELSE IF Blocks (1)

ELSE IF functions much like ELSE and IF

```
IF (X < 0.0) THEN      ! This is tried first
    X = A
ELSE IF (X < 2.0) THEN  ! This second
    X = A + (B-A)*(X-1.0)
ELSE IF (X < 3.0) THEN  ! And this third
    X = B + (C-B)*(X-2.0)
ELSE                   ! This is used if none succeed
    X = C
END IF
```

## Including ELSE IF Blocks (2)

You can have as many **ELSE IF**s as you like  
There is only one **END IF** for the whole block

All **ELSE IF**s must come before any **ELSE**  
Checked in order, and the **first** success is taken

You can omit the **ELSE** in such constructs

**ELSE IF** can be spelled **ELSE IF**

# Named IF Statements (1)

The **IF** can be preceded by **<name>** :

And the **END IF** followed by **<name>** – **note!**

And any **ELSE IF/THEN** and **ELSE** **may** be

```
gnole : IF (X < 0.0) THEN
```

```
    X = A
```

```
ELSE IF (X < 2.0) THEN gnole
```

```
    X = A + (B-A)*(X-1.0)
```

```
ELSE gnole
```

```
    X = C
```

```
END IF gnole
```



## Named IF Statements (2)

The **IF construct name** must match and be distinct  
A great help for checking and clarity

- You should name at least all long **IFs**

If you don't nest **IFs** much, this style is fine

```
gnole : IF (X < 0.0) THEN  
    X = A  
ELSE IF (X < 2.0) THEN  
    X = A + (B-A)*(X-1.0)  
ELSE  
    X = C  
END IF gnole
```

# Block Contents

- Almost any **executable statements** are OK

Both kinds of **IF**, complete **loops** etc.

You may never notice the few restrictions

That applies to all of the **block statements**

**IF**, **DO**, **SELECT** etc.

And all of the **blocks** within an **IF** statement

- Avoid deep levels and very long blocks  
Purely because they will confuse human readers

# Example

```
phasetest: IF (state == 1) THEN
    IF (phase < pi_by_2) THEN
        . . .
    ELSE
        . . .
    END IF
ELSE IF (state == 2) THEN phasetest
    IF (phase > pi) PRINT *, 'A bit odd here'
ELSE phasetest
    IF (phase < pi) THEN
        . . .
    END IF
END IF phasetest
```

# Basic Loops (1)

- A single **loop construct**, with variations  
The basic syntax is:

```
[ loop name : ] DO [ [ , ] loop control ]  
    block  
END DO [ loop name ]
```

**loop name** and **loop control** are optional  
With no **loop control**, it loops indefinitely

**END DO** can be spelled **ENDDO**

The **comma** after **DO** is entirely a matter of taste

## Basic Loops (2)

```
DO      ! Implement the Unix 'yes' command
  PRINT *, 'y'
END DO
```

```
yes: DO
  PRINT *, 'y'
END DO yes
```

The **loop name** must match and be distinct

- You should name at least all long loops

A great help for checking and clarity

Other of its uses are described later

# Indexed Loop Control

The **loop control** has the following form

**<integer variable> = <LWB> , <UPB>**

The **bounds** can be any **integer expressions**

The **variable** starts at the **lower bound**

**A:** If it exceeds the **upper bound**, the loop **exits**

The loop **body** is executed †

The **variable** is **incremented by one**

The loop starts again from **A**

† See later about **EXIT** and **CYCLE**

# Examples

```
DO I = 1 , 3  
    PRINT *, 7*I-3  
END DO
```

Prints 3 lines containing 4, 11 and 18

```
DO I = 3 , 1  
    PRINT *, 7*I-3  
END DO
```

Does nothing

# Using an increment

The general form is

`<var> = <start> , <finish> , <step>`

`<var>` is set to `<start>`, as before

`<var>` is incremented by `<step>`, not `one`

Until it `exceeds` `<finish>` (if `<step>` is `positive`)

Or is `smaller than` `<finish>` (if `<step>` is `negative`)

- The `direction` depends on the `sign` of `<step>`

The loop is `invalid` if `<step>` is `zero`, of course



# Examples

```
DO I = 1 , 20 , 7  
    PRINT *, I  
END DO
```

Prints 3 lines containing 1, 8 and 15

```
DO I = 20 , 1 , 7  
    PRINT *, I  
END DO
```

Does nothing

# Examples

```
DO I = 20 , 1 , -7  
    PRINT *, I  
END DO
```

Prints 3 lines containing 20, 13 and 6

```
DO I = 1 , 20 , -7  
    PRINT *, I  
END DO
```

Does nothing

# Mainly for C Programmers

The **control expressions** are calculated on entry

- Changing their **variables** has no effect
- It is illegal to assign to the **loop variable**

```
DO index = i*j, n**21, k
```

```
    n = 0; k = -1      ! Does not affect the loop
```

```
    index = index+1    ! Is forbidden
```

```
END DO
```

# Loop Control Statements

**EXIT** leaves the innermost loop

**CYCLE** skips to the next iteration

**EXIT/CYCLE name** is for the loop named **name**

These are usually used in single-statement **IFs**

**DO**

**x = read\_number()**

**IF (x < 0.0) EXIT**

**count = count+1; total = total+x**

**IF (x == 0.0) CYCLE**

**...**

**END DO**

# Example

```
INTEGER :: state(right), table(left , right)
FirstMatch = 0
outer: DO i = 1 , right
    IF (state(right) /= OK) CYCLE
    DO j = 1 , left
        IF (found(table(j,i)) THEN
            FirstMatch = i
            EXIT outer
        END IF
    END DO
END DO outer
```

# Warning

What is the **control variable**'s value after **loop exit**?

- It is **undefined** after **normal exit**

Web pages and ignoramuses often say otherwise

It **IS** defined if you leave by **EXIT**

Generally, it is better not to rely on that fact

# WHILE Loop Control

The **loop control** has the following form

**WHILE** ( <logical expression> )

The **expression** is **reevaluated** for each **cycle**

The loop **exits** as soon as it becomes **.FALSE.**

The following are equivalent:

**DO WHILE** ( <logical expression> )

**DO**

**IF** (.NOT. ( <logical expression> )) **EXIT**

# CONTINUE

**CONTINUE** is a statement that does nothing  
It used to be fairly common, but is now rare

Its main use is in **blocks** that do nothing  
**Empty** blocks aren't allowed in Fortran

Otherwise mainly a placeholder for **labels**  
This is **purely** to make the code clearer

But it can be used anywhere a **statement** can



# RETURN and STOP

**RETURN** returns from a **procedure**

- It does **not** return a **result**

How to do that is covered under **procedures**

**STOP** halts the **program** cleanly

- Do **not** spread it throughout your code

Call a **procedure** to tidy up and finish off

# Multi-way IFs

```
IF (expr == val1) THEN
    x = 1.23
ELSE IF (expr >= val2 .AND. expr <= val3) THEN
    CONTINUE
ELSE IF (expr == val4) THEN
    x = x + 4.56
ELSE
    x = 7.89 - x
END IF
```

Very commonly, **expr** is always the same  
And all of the **vals** are **constant expressions**  
**Then** there is another way of coding it

# SELECT CASE (1)

```
PRINT *, 'Happy Birthday'
SELECT CASE (age)
CASE(18)
    PRINT *, 'You can now vote'
CASE(40)
    PRINT *, 'And life begins again'
CASE(60)
    PRINT *, 'And free prescriptions'
CASE(100)
    PRINT *, 'And greetings from the Queen'
CASE DEFAULT
    PRINT *, 'It''s just another birthday'
END SELECT
```

# SELECT CASE (2)

- The **CASE** clauses are **statements**

To put on one line, use '**CASE(18) ; <statement>**'

The values must be **constant expressions**

**INTEGER**, **CHARACTER** or **LOGICAL**

You can specify ranges for the first two

**CASE (-42:42)**      ! -42 to 42 inclusive

**CASE (42:)**      ! 42 or above

**CASE (:42)**      ! Up to and including 42

Be careful with **CHARACTER** ranges

# SELECT CASE (3)

SELECT CASE can be spelled SELECTCASE  
END SELECT can be spelled ENDSELECT

- CASE DEFAULT but NOT CASEDEFAULT

SELECT and CASE can be named, like IF

- It is an error for the ranges to overlap

It is not an error for ranges to be empty

Empty ranges don't overlap with anything

It is not an error for the default to be unreachable

# Labels and GOTO

**Warning:** this area gets seriously religious!

Most **executable statements** can be **labelled**  
**GOTO** **<label>** branches directly to the **label**

In old Fortran, you needed to use a lot of these

- Now, you should almost never use them

If you think you need to, consider redesigning

- **Named loops**, **EXIT** and **CYCLE** are better

# Remaining uses of GOTO

- Useful for branching to clean-up code  
E.g. diagnostics, undoing partial updates etc.  
This is by **FAR** the main remaining use

Fortran does not have any cleaner mechanisms  
E.g. it has no exception handling constructs

- They make a **few** esoteric algorithms clearer  
E.g. certain finite-state machine models  
I have **seen** such code **3–4** times in **40+** years

# Clean-up Code (1)

```
SUBROUTINE Fred
DO . . .
    CALL SUBR (arg1 , arg2 , . . . , argn , ifail)
    IF (ifail /= 0) GOTO 999
END DO
. . . lots more similar code . . .
RETURN

999 SELECT CASE (ifail)
CASE(1) ! Code for ifail = 1
    . . .
CASE(2) ! Code for ifail = 2
    . . .
END SUBROUTINE Fred
```



## Clean-up Code (2)

Many people regard this as better style:

```
SUBROUTINE Fred
DO . . .
    CALL SUBR (arg1 , arg2 , . . . , argn , ifail)
    IF (ifail /= 0) GOTO 999
END DO

999 CONTINUE
SELECT CASE (ifail)
CASE(1) ! Code for ifail = 1
    . . .
END SUBROUTINE Fred
```

# Other Mechanisms

Switches, branches and labels are omitted

They're there in the notes, if you are interested

- You very rarely need to use them, anyway

# Introduction to Modern Fortran

## *Array Concepts*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Array Declarations

Fortran is **the** array-handling language  
Applications like **Matlab** descend from it

You can do almost everything you want to

- Provided that your arrays are **rectangular**  
**Irregular** arrays are possible via **pointers**

- Start by using the simplest features only  
When you need more, check what Fortran has

We will cover the basics and a bit more

# Array Declarations

**Attributes** qualify the **type** in **declarations**  
Immediately following, separated by a **comma**

The **DIMENSION** **attribute** declares arrays  
It has the form **DIMENSION(<dimensions>)**  
Each **<dimension>** is **<lwb>:<upb>**

For example:

```
INTEGER, DIMENSION(0:99) :: table  
REAL, DIMENSION(-10:10, -10:10) :: values
```

# Examples of Declarations

Some examples of array declarations:

```
INTEGER, DIMENSION(0:99) :: arr1, arr2, arr3
INTEGER, DIMENSION(1:12) :: days_in_month
CHARACTER(LEN=10), DIMENSION(1:250) :: names
CHARACTER(LEN=3), DIMENSION(1:12) :: months
REAL, DIMENSION(1:350) :: box_locations
REAL, DIMENSION(-10:10, -10:10) :: pos1, pos2
REAL, DIMENSION(0:5, 1:7, 2:9, 1:4, -5:-2) :: bizarre
```

# Lower Bounds of One

Lower bounds of one (**1:**) can be omitted

```
INTEGER, DIMENSION(12) :: days_in_month  
CHARACTER(LEN=10), DIMENSION(250) :: names  
CHARACTER(LEN=3), DIMENSION(12) :: months  
REAL, DIMENSION(350) :: box_locations  
REAL, DIMENSION(0:5, 7, 2:9, 4, -5:-2) :: bizarre
```

It is entirely a matter of taste whether you do

- C/C++/Python users note **ONE** not **ZERO**

# Alternative Form

The same **base type** but different **bounds**

```
INTEGER :: arr1(0:99), arr2(0:99), arr3(0:99), &  
          days_in_month(1:12)
```

```
REAL :: box_locations(1:350), &  
        pos1(-10:10, -10:10), pos2(-10:10, -10:10), &  
        bizarre(0:5, 1:7, 2:9, 1:4, -5:-2)
```

But this is thoroughly confusing:

```
INTEGER, DIMENSION(0:99) :: arr1, arr2, arr3, &  
          days_in_month(1:12), extra_array, &  
          days_in_leap_year(1:12)
```



# Terminology (1)

```
REAL :: A(0:99), B(3, 6:9, 5)
```

The **rank** is the number of **dimensions**  
**A** has **rank 1** and **B** has **rank 3**

The **bounds** are the upper and lower limits  
**A** has **bounds 0:99** and **B** has **1:3**, **6:9** and **1:5**

A dimension's **extent** is the **UPB-LWB+1**  
**A** has **extent 100** and **B** has **extents 3**, **4** and **5**

# Terminology (2)

```
REAL :: A(0:99), B(3, 6:9, 5)
```

The **size** is the total number of **elements**  
**A** has **size** 100 and **B** has **size** 60

The **shape** is its **rank** and **extents**  
**A** has **shape** (100) and **B** has **shape** (3,4,5)

Arrays are **conformable** if they share a **shape**

- The **bounds** do not have to be the same

# Array Element References

An array **index** can be any **integer** expression  
E.g. **months(J)**, selects the **J**th month

```
INTEGER, DIMENSION(-50:50) :: mark
DO I = -50, 50
    mark(I) = 2*I
END DO
```

Sets **mark** to **-100, -98, ..., 98, 100**

# Index Expressions

```
INTEGER, DIMENSION(1:80) :: series
DO K = 1, 40
    series(2*K) = 2*K-1
    series(2*K-1) = 2*K
END DO
```

Sets the **even elements** to the **odd indices**  
And vice versa

You can go completely overboard, too

```
series(int(1.0+80.0*cos(123.456))) = 42
```

# Example of Arrays – Sorting

Sort a list of numbers into ascending order  
The top-level algorithm is:

1. Read the numbers and store them in an array.
2. Sort them into ascending order of magnitude.
3. Print them out in sorted order.

# Selection Sort

This is **NOT** how to write a general sort

It takes  $O(N^2)$  time – compared to  $O(N \log(N))$

For each location **J** from **1** to **N-1**

    For each location **K** from **J+1** to **N**

        If the value at **J** exceeds that at **K**

            Then swap them

    End of loop

End of loop

# Selection Sort (1)

```
PROGRAM sort10
  INTEGER, DIMENSION(1:10) :: nums
  INTEGER :: temp, J, K
! --- Read in the data
  PRINT *, 'Type ten integers each on a new line'
  DO J = 1, 10
    READ *, nums(J)
  END DO
! --- Sort the numbers into ascending order of magnitude
  . . .
! --- Write out the sorted list
  DO J = 1, 10
    PRINT *, 'Rank ', J, ' Value is ', nums(J)
  END DO
END PROGRAM sort10
```

# Selection Sort (2)

! --- Sort the numbers into ascending order of magnitude

L1: DO J = 1, 9

L2: DO K = J+1, 10

IF(nums(J) > nums(K)) THEN

temp = nums(K)

nums(K) = nums(J)

nums(J) = temp

END IF

END DO L2

END DO L1



# Valid Array Bounds

The **bounds** can be any **constant expressions**

There are two ways to use **run-time** bounds

- **ALLOCATABLE** arrays – see later
- When allocating them in **procedures**

We will discuss the following under **procedures**

```
SUBROUTINE workspace (size)
```

```
  INTEGER :: size
```

```
  REAL, DIMENSION(1:size*(size+1)) :: array
```

```
  . . .
```

# Using Arrays as Objects (1)

Arrays can be handled as **compound objects**  
**Sections** allow access as groups of elements  
There are a large number of **intrinsic procedures**

Simple use handles all elements “in parallel”

- **Scalar** values are expanded as needed

Set all elements of an array to a single value

```
INTEGER, DIMENSION(1:50) :: mark  
mark = 0
```

# Using Arrays as Objects (2)

You can use **whole arrays** as simple **variables**  
Provided that they are all **conformable**

```
REAL, DIMENSION(1:200) :: arr1, arr2
```

```
• • •
```

```
arr1 = arr2+1.23*exp(arr1/4.56)
```

- I really do mean “**as simple variables**”

The **RHS** and any **LHS** indices are evaluated  
And **then** the **RHS** is assigned to the **LHS**

# Array Sections

Array **sections** create an aliased subarray  
It is a **simple variable** with a **value**

```
INTEGER :: arr1(1:100), arr2(1:50), arr3(1:100)
```


```
arr1(1:63) = 5 ;   arr1(64:100) = 7  
arr2 = arr1(1:50)+arr3(51:100)
```

- Even this is legal, but forces a **copy**

```
arr1(26:75) = arr1(1:50)+arr1(51:100)
```

# Array Sections

**A(1:6,1:8)**



# Short Form

Existing array bounds may be omitted  
Especially useful for multidimensional arrays

If we have `REAL, DIMENSION(1:6, 1:8) :: A`  
`A(3:, :4)` is the same as `A(3:6, 1:4)`  
`A`, `A(:, :)` and `A(1:6, 1:8)` are all the same

`A(6, :)` is the 6th row as a 1-D vector

`A(:, 3)` is the 3rd column as a 1-D vector

`A(6:6, :)` is the 6th row as a  $1 \times 8$  matrix

`A(:, 3:3)` is the 3rd column as a  $6 \times 1$  matrix

# Conformability of Sections

The **conformability** rule applies to sections, too

```
REAL :: A(1:6, 1:8), B(0:3, -5:5), C(0:10)
```

```
A(2:5, 1:7) = B(:, -3:3)    ! both have shape (4, 7)
```

```
A(4, 2:5) = B(:, 0) + C(7:)  ! all have shape (4)
```

```
C(:) = B(2, :)    ! both have shape (11)
```

But these would be illegal

```
A(1:5, 1:7) = B(:, -3:3)    ! shapes (5,7) and (4,7)
```

```
A(1:1, 1:3) = B(1, 1:3)    ! shapes (1,3) and (3)
```

# Sections with Strides

Array sections need not be **contiguous**

Any **uniform progression** is allowed

This is **exactly** like a more compact **DO**-loop

Negative strides are allowed, too

```
INTEGER :: arr1(1:100), arr2(1:50), arr3(1:50)
```

```
arr1(1:100:2) = arr2    ! Sets every odd element
```


```
arr1(100:1:-2) = arr3   ! Even elements, reversed
```

```
arr1 = arr1(100:1:-1)   ! Reverses the order of arr1
```



# Strided Sections

**A(1:6,1:8)**



# Array Bounds

Subscripts/sections must be within bounds

The following are **invalid** (undefined behaviour)

```
REAL :: A(1:6, 1:8), B(0:3, -5:5), C(0:10)
```

```
A(2:5, 1:7) = B(:, -6:3)
```

```
A(7, 2:5) = B(:, 0)
```

```
C(:11) = B(2, :)
```

**NAG** will usually check; most others won't

Errors lead to overwriting etc. and **CHAOS**

Even **NAG** may not check all **old-style** Fortran

# Elemental Operations

We have seen **operations** and **intrinsic functions**  
Most built-in operators/functions are **elemental**  
They act **element-by-element** on arrays

```
REAL, DIMENSION(1:200) :: arr1, arr2, arr3  
arr1 = arr2+1.23*exp(arr3/4.56)
```

Comparisons and logical operations, too

```
REAL, DIMENSION(1:200) :: arr1, arr2, arr3  
LOGICAL, DIMENSION(1:200) :: flags  
flags = (arr1 > exp(arr2) .OR. arr3 < 0.0)
```

# Array Intrinsic Functions (1)

There are over 20 useful **intrinsic procedures**  
They can save a lot of coding and debugging

SIZE(x [, n])	! The size of x (an integer scalar)
SHAPE(x)	! The shape of x (an integer vector)

LBOUND(x [, n])	! The lower bound of x
UBOUND(x [, n])	! The upper bound of x

If **n** is present, down that dimension only  
And the result is is an **integer scalar**  
Otherwise the result is is an **integer vector**

# Array Intrinsic Functions (2)

MINVAL(x)	! The minimum of all elements of x
MAXVAL(x)	! The maximum of all elements of x

These return a **scalar** of the same **type** as **x**

MINLOC(x)	! The indices of the minimum
MAXLOC(x)	! The indices of the maximum

These return an **integer vector**, just like **SHAPE**

# Array Intrinsic Functions (3)

SUM(x [, n])	! The sum of all elements of x
PRODUCT(x [, n])	! The product of all elements of x

If **n** is present, down that dimension only

TRANSPOSE(x)	! The transposition of
DOT_PRODUCT(x, y)	! The dot product of x and y
MATMUL(x, y)	! 1- and 2-D matrix multiplication

# Reminder

TRANPOSE(X) means  $X_{ij} \Rightarrow X_{ji}$

It must have **two** dimensions, but needn't be **square**

DOT\_PRODUCT(X, Y) means  $\sum_i X_i \cdot Y_i \Rightarrow Z$

Two vectors, both of the same length and type

MATMUL(X, Y) means  $\sum_k X_{ik} \cdot Y_{kj} \Rightarrow Z_{ij}$

**Second** dimension of **X** must match the **first** of **Y**

The matrices need not be the same **shape**

Either of **X** or **Y** may be a **vector** in **MATMUL**

# Array Intrinsic Functions (4)

These also have some features not mentioned

There are more (especially for [reshaping](#))

There are ones for [array masking](#) (see later)

Look at the references for the details



# Warning

It's not specified **how** results are calculated  
All of the following can be different:

- Calling the **intrinsic function**
- The **obvious code** on array elements
- The **numerically best** way to do it
- The **fastest** way to do it

All of them are calculate the same **formula**  
But the **result** may be slightly different

- If this starts to matter, consult an expert

# Array Element Order (1)

This is also called “storage order”

Traditional term is “column-major order”

But Fortran arrays are not laid out in columns!

Much clearer: “first index varies fastest”

```
REAL :: A(1:3, 1:4)
```

The elements of **A** are stored in the order

```
A(1,1), A(2,1), A(3,1), A(1,2), A(2,2), A(3,2), A(1,3),  
A(2,3), A(3,3), A(1,4), A(2,4), A(3,4)
```

# Array Element Order (2)

Opposite to C, Matlab, Mathematica etc.

You don't often need to know the storage order  
**Three** important cases where you do:

- I/O of arrays, especially unformatted
- Array constructors and array constants
- Optimisation (caching and locality)

There are more cases in **old-style** Fortran  
Avoid that, and you need not learn them

# Simple I/O of Arrays (1)

**Arrays** and **sections** can be included in I/O  
These are expanded in **array element order**

```
REAL, DIMENSION(3, 2) :: oxo  
READ *, oxo
```

This is **exactly** equivalent to:

```
REAL, DIMENSION(3, 2) :: oxo  
READ *, oxo(1, 1), oxo(2, 1), oxo(3, 1), &  
        oxo(1, 2), oxo(2, 2), oxo(3, 2)
```

# Simple I/O of Arrays (2)

Array sections can also be used

```
REAL, DIMENSION(100) :: nums  
READ *, nums(30:50)
```

```
REAL, DIMENSION(3, 3) :: oxo  
READ *, oxo(:, 3), oxo(3:1:-1, 1)
```

The last statement is equivalent to

```
READ *, oxo(1, 3), oxo(2, 3), oxo(3, 3),    &  
        oxo(3, 1), oxo(2, 1), oxo(1, 1)
```

# Array Constructors (1)

An **array constructor** creates a temporary array

- Commonly used for assigning array values

```
INTEGER :: marks(1:6)  
marks = (/ 10, 25, 32, 54, 54, 60 /)
```

Constructs an array with elements

10, 25, 32, 54, 54, 60

And then copies that array into **marks**

A good compiler will optimise that!

# Array Constructors (2)

- Variable expressions are OK in constructors

(/ x, 2.0\*y, SIN(t\*w/3.0),... etc. /)

They can be used anywhere an array can be  
Except where you might assign to them!

- All expressions must be the same type
- This has been relaxed in Fortran 2003

# Array Constructors (3)

Arrays can be used in the value list  
They are flattened into array element order

Implied DO-loops (as in I/O) allow sequences

If  $n$  has the value 7

$(/ 0.0, (k/10.0, k = 2, n), 1.0 /)$

Is equivalent to:

$(/ 0.0, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 1.0 /)$



# Constants and Initialisation (1)

Array constructors are very useful for this  
All elements must be **constant expressions**  
I.e. ones that can be evaluated at compile time

For **rank one** arrays, just use a constructor

```
REAL, PARAMETER :: a(1:3) = (/ 1.23, 4.56, 7.89 /)
REAL, PARAMETER :: b(3) = exp( (/ 1.2, 3.4, 5.6 /) )
```

But **NOT**:

```
REAL, PARAMETER :: arr(1:3) = &
    myfunc ( (/ 1.2, 3.4, 5.6 /) )
```

# Constants and Initialisation (2)

Other types can be initialised in the same way

```
CHARACTER(LEN=4), DIMENSION(1:5) :: names = &  
(/ 'Fred', 'Joe', 'Bill', 'Bert', 'Alf' /)
```

Constant expressions are allowed

```
INTEGER, PARAMETER :: N = 3, M = 6, P = 12  
INTEGER :: arr(1:3) = (/ N, (M/N), (P/N) /)  
REAL :: arr(1:3) = (/ 1.0, exp(1.0), exp(2.0) /)
```

But **NOT**:

```
REAL :: arr(1:3) = (/ 1.0, myfunc(1.0), myfunc(2.0) /)
```

# Multiple Dimensions

Constructors cannot be nested – e.g. **NOT**:

```
REAL, DIMENSION(3, 4) :: array = &  
    (/ (/ 1.1, 2.1, 3.1 /), (/ 1.2, 2.2, 3.2 /), &  
        (/ 1.3, 2.3, 3.3 /), (/ 1.4, 2.4, 3.4 /) /)
```

They construct only **rank one** arrays

- Construct higher ranks using **RESHAPE**  
This is covered in the extra slides on arrays

# Allocatable Arrays (1)

Arrays can be declared with an **unknown shape**  
Attempting to use them in that state will fail

```
INTEGER, DIMENSION(:), ALLOCATABLE :: counts  
REAL, DIMENSION(:, :, :), ALLOCATABLE :: values
```

They become defined when space is allocated

```
ALLOCATE (counts(1:1000000))  
ALLOCATE (value(0:N, -5:5, M:2*N+1))
```

# Allocatable Arrays (2)

Failure will terminate the program

You can trap most allocation failures

```
INTEGER :: istat  
ALLOCATE (arr(0:100, -5:5, 7:14), STAT=istat)  
IF (istat /= 0) THEN  
    . . .  
END IF
```

Arrays can be deallocated using

```
DEALLOCATE (nums)
```

There are more features in [Fortran 2003](#)

# Example

```
INTEGER, DIMENSION(:), ALLOCATABLE :: counts
INTEGER :: size, code
! --- Ask the user how many counts he has
PRINT *, 'Type in the number of counts'
READ *, size
! --- Allocate memory for the array
ALLOCATE (counts(1:size), STAT=code)
IF (code /= 0) THEN
    . . .
END IF
```

# Allocation and Fortran 95

Fortran 95 constrained **ALLOCATABLE** objects  
Cannot be arguments, results or in derived types  
I.e. local to procedures or in modules only

Fortran 2003 allows them almost everywhere  
Almost all compilers already include those features  
You may come across **POINTER** in old code  
It can usually be replaced by **ALLOCATABLE**

Ask if you hit problems and want to check

# Testing Allocation

Can test if an **ALLOCATABLE** object is **allocated**

The **ALLOCATED** function returns **LOGICAL**:

**INTEGER, DIMENSION(:), ALLOCATABLE :: counts**

• • •  
**IF (ALLOCATED(counts)) THEN**

• • •

Generally, that is needed for **advanced use** only



# Allocatable CHARACTER

Remember **CHARACTER** is really a **string**  
Not an **array** of single characters, but a bit like one

You can use a colon (:) for a **length**  
Provided that the variable is **allocatable**

This makes a copy of the text on an input line  
It is also a **Fortran 2003** feature

```
CHARACTER(LEN=100) :: line  
CHARACTER(LEN=:), ALLOCATABLE :: message  
ALLOCATE (message, SOURCE=TRIM(line))
```

# Reminder

The above is all many programmers need  
There is a lot more, but skip it for now

At this point, let's see a real example  
**Cholesky decomposition** following **LAPACK**  
With all error checking omitted, for clarity

It isn't pretty, but it is like the mathematics

- And that **really** helps to reduce errors  
E.g. coding up a published algorithm

# Cholesky Decomposition

To solve  $A = LL^T$ , in tensor notation:

$$L_{jj} = \sqrt{A_{jj} - \sum_{k=1}^{j-1} L_{jk}^2}$$

$$\forall i > j, L_{ij} = (A_{ij} - \sum_{k=1}^{j-1} L_{ik} L_{jk}) / L_{jj}$$

Most of the Web uses  $i$  and  $j$  the other way round

# Cholesky Decomposition

```
SUBROUTINE CHOLESKY ( A )  
  IMPLICIT NONE  
  INTEGER :: J, N  
  REAL :: A (:, :)  
  N = UBOUND (A, 1)  
  DO J = 1, N  
    A(J, J) = SQRT ( A(J, J) - &  
      DOT_PRODUCT ( A(J, :J-1), A(J, :J-1) ) )  
    IF ( J < N ) &  
      A(J+1:, J) = ( A(J+1:, J) - &  
        MATMUL ( A(J+1:, :J-1), A(J, :J-1) ) ) / A(J, J)  
  END DO  
END SUBROUTINE CHOLESKY
```

# Other Important Features

These have been omitted for simplicity

There are extra slides giving an overview

- Constructing **higher rank** array constants
- Using **integer vectors** as **indices**
- **Masked assignment** and **WHERE**
- **Memory locality** and **performance**
- Avoiding unnecessary **array copying**

# Introduction to Modern Fortran

## *Procedures*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Sub-Dividing The Problem

- Most programs are **thousands** of lines  
Few people can grasp all the details
- You often use similar code in several places
- You often want to test parts of the code
- Designs often break up naturally into steps

Hence, all sane programmers use **procedures**

# What Fortran Provides

There must be a single **main program**

There are **subroutines** and **functions**

All are collectively called **procedures**

A **subroutine** is some **out-of-line** code

There are very few restrictions on what it can do

It is always called exactly where it is coded

A **function**'s purpose is to return a **result**

There are some restrictions on what it can do

It is called only when its **result** is needed



# Example – Cholesky (1)

We saw this when considering **arrays**  
It is a very typical, simple **subroutine**

```
SUBROUTINE CHOLESKY (A)
  IMPLICIT NONE
  INTEGER :: J, N
  REAL :: A(:, :), X
  N = UBOUND(A, 1)
  DO J = 1, N
    ...
  END DO
END SUBROUTINE CHOLESKY
```

# Example – Cholesky (2)

And this is how it is **called**

```
PROGRAM MAIN
  IMPLICIT NONE
  REAL, DIMENSION(5, 5) :: A = 0.0
  REAL, DIMENSION(5) :: Z
  . . .
  CALL CHOLESKY (A)
  . . .
END PROGRAM MAIN
```

We shall see how to declare it later

# Example – Variance

```
FUNCTION Variance (Array)
  IMPLICIT NONE
  REAL :: Variance, X
  REAL, INTENT(IN), DIMENSION(:) :: Array
  X = SUM(Array)/SIZE(Array)
  Variance = SUM((Array-X)**2)/SIZE(Array)
END FUNCTION Variance
```

```
REAL, DIMENSION(1000) :: data
```

```
...
```

```
Z = Variance(data)
```

We shall see how to declare it later

# Example – Sorting (1)

This was the harness of the selection sort  
Replace the actual sorting code by a call

```
PROGRAM sort10
  IMPLICIT NONE
  INTEGER, DIMENSION(1:10) :: nums
  . . .
  ! --- Sort the numbers into ascending order of magnitude
  CALL SORTIT (nums)
  ! --- Write out the sorted list
  . . .
END PROGRAM sort10
```

## Example – Sorting (2)

```
SUBROUTINE SORTIT (array)
  IMPLICIT NONE
  INTEGER :: temp, array(:), J, K
L1:   DO J = 1, UBOUND(array,1)-1
L2:     DO K = J+1, UBOUND(array,1)
        IF(array(J) > array(K)) THEN
            temp = array(K)
            array(K) = array(J)
            array(J) = temp
        END IF
      END DO L2
    END DO L1
END SUBROUTINE SORTIT
```

# CALL Statement

The **CALL** statement **evaluates** its arguments  
The following is an **over-simplified** description

- **Variables** and **array sections** define memory
- **Expressions** are stored in a **hidden variable**

It then transfers control to the **subroutine**  
Passing the locations of the **actual arguments**

Upon **return**, the next statement is executed

# SUBROUTINE Statement

Declares the **procedure** and its **arguments**  
These are called **dummy arguments** in Fortran

The subroutine's **interface** is defined by:

- The **SUBROUTINE** statement itself
- The **declarations** of its **dummy arguments**
- And anything that those use (see later)

```
SUBROUTINE SORTIT (array)
```

```
INTEGER :: [ temp, ] array(:) [ , J, K ]
```

# Subroutines With No Arguments

You aren't required to have any **arguments**

You can omit the **parentheses** if you prefer

Preferably either **do** or **don't**, but you can mix uses

```
SUBROUTINE Joe ( )
```

```
SUBROUTINE Joe
```

```
CALL Joe ( )
```

```
CALL Joe
```



# Statement Order

A **SUBROUTINE** statement starts a **subroutine**

Any **USE** statements must come next

Then **IMPLICIT NONE**

Then the rest of the **declarations**

Then the **executable statements**

It ends at an **END SUBROUTINE** statement

**PROGRAM** and **FUNCTION** are similar

There are other rules, but you may ignore them

# Dummy Arguments

- Their **names** exist only in the **procedure**  
They are declared much like **local variables**

Any **actual argument** names are irrelevant  
Or any other names outside the **procedure**

- The **dummy arguments** are associated  
with the **actual arguments**

Think of **association** as a bit like **aliasing**

# Argument Matching

Dummy and actual argument lists must match

The number of arguments must be the same

Each argument must match in type and rank

That can be relaxed in several ways

See under advanced use of procedures

We shall come back to array arguments shortly

Most of the complexities involve them

This is for compatibility with old standards

# Functions

Often the required result is a single value  
It is easier to write a **FUNCTION** procedure

E.g. to find the largest of three values:

- Find the largest of the first and second
- Find the largest of that and the third

Yes, I know that the **MAX** function does this!

The **function name** defines a **local variable**

- Its value on **return** is the result returned

The **RETURN** statement does not take a value

# Example (1)

```
FUNCTION largest_of (first, second, third)
  IMPLICIT NONE
  INTEGER :: largest_of
  INTEGER :: first, second, third
  IF (first > second) THEN
    largest_of = first
  ELSE
    largest_of = second
  END IF
  IF (third > largest_of) largest_of = third
END FUNCTION largest_of
```

## Example (2)

```
INTEGER :: trial1, trial2 ,trial3, total, count
total = 0 ;   count = 0
DO
    PRINT *, 'Type three trial values:'
    READ *, trial1, trial2, trial3
    IF (MIN(trial1, trial2, trial3) < 0) EXIT
        count = count + 1
        total = total + &
            largest_of(trial1, trial2, trial3)
END DO
PRINT *, 'Number of trial sets = ', count, &
    ' Total of best of 3 = ',total
```

# Functions With No Arguments

You aren't required to have any arguments

You must **not** omit the parentheses

```
FUNCTION Fred ( )  
  INTEGER :: Fred
```

```
X = 1.23 * Fred ( )  
CALL Alf ( Fred ( ) )
```

In the following, **Fred** is a procedure argument

```
CALL Alf ( Fred )
```

# Internal Procedures (1)

Procedures can contain internal procedures  
These can be SUBROUTINEs and FUNCTIONs  
The statement order is as follows:

PROGRAM, SUBROUTINE or FUNCTION

All of the code of the actual procedure

CONTAINS

Any number of internal procedures

END PROGRAM, SUBROUTINE or FUNCTION

- Internal procedures may not themselves contain internal procedures



# Internal Procedures (2)

- **Warning:** that order takes some getting used to

The **procedure** can use the **internal procedures**  
And **one of them** can call **any other**

Most useful for small, private **auxiliary** ones  
You can include any number of internal procedures

- They are visible only in the **outer procedure**  
Won't **clash** with the **same name** elsewhere

# Internal Procedures (3)

```
PROGRAM main
  REAL, DIMENSION(10) :: vector
  PRINT *, 'Type 10 values'
  READ *, vector
  PRINT *, 'Variance = ', Variance(vector)
CONTAINS
  FUNCTION Variance (Array)
    REAL :: Variance, X
    REAL, INTENT(IN), DIMENSION(:) :: Array
    X = SUM(Array)/SIZE(Array)
    Variance = SUM((Array-X)**2)/SIZE(Array)
  END FUNCTION Variance
END PROGRAM main
```

# Name Inheritance (1)

Everything accessible in the **enclosing procedure**  
can also be used in the **internal procedure**

This includes **all** of the local declarations  
And anything imported by **USE** (covered later)

Internal procedures need only a few arguments  
Just the things that vary between calls  
Everything else can be used directly

# Name Inheritance (2)

A **local name** takes precedence

```
PROGRAM main
  REAL :: temp = 1.23
  CALL pete (4.56)
CONTAINS
  SUBROUTINE pete (temp)
    PRINT *, temp
  END SUBROUTINE pete
END PROGRAM main
```

Will print **4.56**, not **1.23**

Avoid doing this – it's very confusing

# Using Procedures

Use this technique for solving test problems

- It is one of the **best techniques** for real code

There is **another**, equally good one, under **modules**

And there are yet others that you may need to use

# INTENT (1)

You can make arguments **read-only**

```
SUBROUTINE Summarise (array, size)
  INTEGER, INTENT(IN) :: size
  REAL, DIMENSION(size) :: array
```

That will prevent you writing to it by accident  
Or calling another procedure that does that  
It may also help the compiler to optimise

- **Strongly** recommended for **read-only** args

## INTENT (2)

You can also make them **write-only**  
Less useful, but still very worthwhile

```
SUBROUTINE Init (array, value)
  IMPLICIT NONE
  REAL, DIMENSION(:), INTENT(OUT) :: array
  REAL, INTENT(IN) :: value
  array = value
END SUBROUTINE Init
```

As useful for optimisation as **INTENT(IN)**

# INTENT (3)

The default is effectively **INTENT(INOUT)**

- But specifying **INTENT(INOUT)** is useful  
It will trap the following nasty error

```
SUBROUTINE Munge (value)
    REAL, INTENT(INOUT) :: value
    value = 100.0*value
    PRINT *, value
END SUBROUTINE Munge

CALL Munge(1.23)
```



# Example

```
SUBROUTINE expsum(n, k, x, sum)
  IMPLICIT NONE
  INTEGER, INTENT(IN) :: n
  REAL, INTENT(IN) :: k, x
  REAL, INTENT(OUT) :: sum
  INTEGER :: i
  sum = 0.0
  DO i = 1, n
    sum = sum + exp(-i*k*x)
  END DO
END SUBROUTINE expsum
```

# Aliasing

Two arguments may **overlap** only if read-only  
Also applies to **arguments** and **global data**

- If **either** is updated, weird things happen

Fortran doesn't have any way to trap that  
Nor do any other current languages – sorry

Use of **INTENT(IN)** will stop it in many cases

- Be careful when using **array arguments**  
Including using **array elements** as **arguments**

# PURE Functions

You can declare a **function** to be **PURE**

All **data arguments** must specify **INTENT(IN)**

It must not modify any **global data**

It must not do I/O (except with **internal files**)

It must call only **PURE** procedures

Some restrictions on more advanced features

Generally overkill – but good practice

Most **built-in** procedures are **PURE**

# Example

This is the cleanest way to define a function

```
PURE FUNCTION Variance (Array)
  IMPLICIT NONE
  REAL :: Variance, X
  REAL, INTENT(IN), DIMENSION(:) :: Array
  X = SUM(Array)/SIZE(Array)
  Variance = SUM((Array-X)**2)/SIZE(Array)
END FUNCTION Variance
```

Most safety, and best possible optimisation

# ELEMENTAL Functions

Functions can be declared as **ELEMENTAL**  
Like **PURE**, but **arguments** must be **scalar**

You can use them on **arrays** and in **WHERE**  
They apply to each **element**, like **built-in SIN**

```
ELEMENTAL FUNCTION Scale (arg1, arg2)
    REAL, INTENT(IN) :: arg1, arg2
    Scale = arg1/sqrt(arg1**2+arg2**2)
END FUNCTION Scale
```

```
REAL, DIMENSION(100) :: arr1, arr2, array
array = Scale(arr1, arr2)
```

# Keyword Arguments (1)

```
SUBROUTINE AXIS (X0, Y0, Length, Min, Max, Intervals)
    REAL, INTENT(IN) :: X0, Y0, Length, Min, Max
    INTEGER, INTENT(IN) :: Intervals
END SUBROUTINE AXIS
```

```
CALL AXIS(0.0, 0.0, 100.0, 0.1, 1.0, 10)
```

- Error prone to write and unclear to read

And it can be a lot worse than that!

# Keyword Arguments (2)

Dummy arg. names can be used as keywords  
You don't have to remember their order

```
SUBROUTINE AXIS (X0, Y0, Length, Min, Max, Intervals)
```

```
...
```

```
CALL AXIS(Intervals=10, Length=100.0, &  
          Min=0.1, Max=1.0, X0=0.0, Y0=0.0)
```

- The argument order now doesn't matter  
The keywords identify the dummy arguments

# Keyword Arguments (3)

**Keywords** arguments can follow **positional**  
The following is allowed

```
SUBROUTINE AXIS (X0, Y0, Length, Min, Max, Intervals)
```

```
...
```

```
CALL AXIS(0.0, 0.0, Intervals=10, Length=100.0, &  
          Min=0.1, Max=1.0)
```

- Remember that the **best** code is the **clearest**  
Use whichever convention feels most natural



# Keyword Reminder

Keywords are **not** names in the calling procedure  
They are used **only** to map to dummy arguments  
The following works, but is somewhat confusing

```
SUBROUTINE Nuts (X, Y, Z)
    REAL, DIMENSION(:) :: X
    INTEGER :: Y, Z
END SUBROUTINE Nuts
```

```
INTEGER :: X
REAL, DIMENSION(100) :: Y, Z
CALL Nuts (Y=X, Z=1, X=Y)
```

# Hiatus

That is most of the basics of **procedures**  
Except for **arrays** and **CHARACTER**

Now might be a good time to do some examples  
The first few questions cover the material so far

# Assumed Shape Arrays (1)

- The best way to declare **array arguments**  
You must declare **procedures** as above
- Specify all **bounds** as simply a colon (':')  
The **rank** must match the **actual argument**  
The **lower bounds** default to **one** (1)  
The **upper bounds** are taken from the **extents**

REAL, DIMENSION(:) :: vector

REAL, DIMENSION(:, :) :: matrix

REAL, DIMENSION(:, :, :) :: tensor

# Example

```
SUBROUTINE Peculiar (vector, matrix)
  REAL, DIMENSION(:), INTENT(INOUT) :: vector
  REAL, DIMENSION(:, :), INTENT(IN) :: matrix
  . . .
END SUBROUTINE Peculiar
```

```
REAL, DIMENSION(20:1000), :: one
REAL, DIMENSION(-5:100, -5:100) :: two
CALL Peculiar (one(101:160), two(21:, 26:75) )
```

vector will be **DIMENSION(1:60)**

matrix will be **DIMENSION(1:80, 1:50)**

# Assumed Shape Arrays (2)

Query functions were described earlier

**SIZE**, **SHAPE**, **LBOUND** and **UBOUND**

So you can write completely **generic** procedures

```
SUBROUTINE Init (matrix, scale)
  REAL, DIMENSION(:, :), INTENT(OUT) :: matrix
  INTEGER, INTENT(IN) :: scale
  DO N = 1, UBOUND(matrix,2)
    DO M = 1, UBOUND(matrix,1)
      matrix(M, N) = scale*M + N
    END DO
  END DO
END SUBROUTINE Init
```

# Cholesky Decomposition

```
SUBROUTINE CHOLESKY(A)
  IMPLICIT NONE
  INTEGER :: J, N
  REAL, INTENT(INOUT) :: A(:, :), X
  N = UBOUND(A, 1)
  IF (N < 1 .OR. UBOUND(A, 2) /= N)
    CALL Error("Invalid array passed to CHOLESKY")
  DO J = 1, N
    . . .
  END DO
END SUBROUTINE CHOLESKY
```

Now I have added appropriate checking

# Setting Lower Bounds

Even when using **assumed shape arrays**  
you can set any **lower bounds** you want

- You do that in the **called procedure**

```
SUBROUTINE Orrible (vector, matrix, n)
    REAL, DIMENSION(2*n+1:) :: vector
    REAL, DIMENSION(0:, 0:) :: matrix
    . . .
END SUBROUTINE Orrible
```

# Warning

**Argument overlap** will not be detected  
Not even for **assumed shape** arrays

- A common cause of obscure errors

No other language does much better



# Explicit Array Bounds

In **procedures**, they are more flexible  
Any reasonable **integer expression** is allowed

Essentially, you can use any ordinary formula  
Using only **constants** and **integer variables**  
Few programmers will ever hit the restrictions

The most common use is for **workspace**  
But it applies to all **array declarations**

# Automatic Arrays (1)

Local arrays with run-time bounds are called automatic arrays

Bounds may be taken from an argument  
Or a constant or variable in a module

```
SUBROUTINE aardvark (size)
USE sizemod    ! This defines worksize
INTEGER, INTENT(IN) :: size

REAL, DIMENSION(1:worksize) :: array_1
REAL, DIMENSION(1:size*(size+1)) :: array_2
```

## Automatic Arrays (2)

Another very common use is a ‘**shadow**’ array  
i.e. one the same **shape** as an **argument**

```
SUBROUTINE pard (matrix)  
  REAL, DIMENSION(:, :) :: matrix
```

```
  REAL, DIMENSION(UBOUND(matrix, 1), &  
                  UBOUND(matrix, 2)) :: &  
    matrix_2, matrix_3
```

And so on – **automatic arrays** are very flexible

# Explicit Shape Array Args (1)

We cover these because of their importance  
They were the only mechanism in Fortran 77

- But, generally, they should be avoided

In this form, all bounds are explicit  
They are declared just like automatic arrays  
The dummy should match the actual argument  
Making an error will usually cause chaos

- Only the very simplest uses are covered
- There are more details in the extra slides

# Explicit Shape Array Args (2)

You can use **constants**

```
SUBROUTINE Orace (matrix, array)
  INTEGER, PARAMETER :: M = 5, N = 10
  REAL, DIMENSION(1:M, 1:N) :: matrix
  REAL, DIMENSION(1000) :: array
  . . .
END SUBROUTINE Orace
```

```
INTEGER, PARAMETER :: M = 5, N = 10
REAL, DIMENSION(1:M, 1:N) :: table
REAL, DIMENSION(1000) :: workspace
CALL Orace(table, workspace)
```

# Explicit Shape Array Args (3)

It is common to pass the **bounds** as **arguments**

```
SUBROUTINE Weeble (matrix, m, n)
  INTEGER, INTENT(IN) :: m, n
  REAL, DIMENSION(1:m, 1:n) :: matrix
  . . .
END SUBROUTINE Weeble
```

You can use **expressions**, of course

- But it is not really recommended

Purely on the grounds of human confusion

# Explicit Shape Array Args (4)

You can define the **bounds** in a **module**  
Either as a **constant** or in a **variable**

```
SUBROUTINE Wobble (matrix)
    USE sizemod    ! This defines m and n
    REAL, DIMENSION(1:m, 1:n) :: matrix
    . . .
END SUBROUTINE Weeble
```

- The same remarks about **expressions** apply

# Assumed Size Array Args

The last upper bound can be \*

I.e. **unknown**, but assumed to be large enough

```
SUBROUTINE Weeble (matrix, n)
    REAL, DIMENSION(n, *) :: matrix
    . . .
END SUBROUTINE Weeble
```

- You will see this, but generally avoid it  
It makes it very hard to locate **bounds errors**  
It also implies several restrictions



# Warnings

The **size** of the **dummy array** must not exceed the **size** of the **actual array** argument

- Compilers will rarely detect this error

There are also some performance problems when passing **assumed shape** and **array sections** to **explicit shape** or **assumed size** dummies

That is in the **advanced** slides on procedures  
Sorry – but it's complicated to explain

# Example (1)

We have a subroutine with an interface like:

```
SUBROUTINE Normalise (array, size)
INTEGER, INTENT(IN) :: size
REAL, DIMENSION(size) :: array
```

The following calls are correct:

```
REAL, DIMENSION(1:10) :: data
```

```
CALL Normalise (data, 10)
```

```
CALL Normalise (data(2:5), SIZE(data(2:5)))
```

```
CALL Normalise (data, 7)
```

## Example (2)

```
SUBROUTINE Normalise (array, size)
INTEGER, INTENT(IN) :: size
REAL, DIMENSION(size) :: array
```

The following calls are **not** correct:

```
INTEGER, DIMENSION(1:10) :: indices
REAL :: var, data(10)
```

```
CALL Normalise (indices, 10)  ! wrong base type
CALL Normalise (var, 1)      ! not an array
CALL Normalise (data, 10.0)  ! wrong type
CALL Normalise (data, 20)    ! dummy array too big
```

# Character Arguments

Few scientists do anything very fancy with these  
See the advanced foils for anything like that

People often use a **constant** length  
You can specify this as a **digit string**

Or define it using **PARAMETER**  
That is best done in a module

Or define it as an **assumed length** argument

# Explicit Length Character (1)

The **dummy** should match the **actual argument**  
You are likely to get confused if it doesn't

```
SUBROUTINE sorter (list)
    CHARACTER(LEN=8), DIMENSION(:) :: list
    . . .
END SUBROUTINE sorter

CHARACTER(LEN=8) :: data(1000)
. . .
CALL sorter(data)
```

# Explicit Length Character (2)

```
MODULE Constants
```

```
    INTEGER, PARAMETER :: charlen = 8
```

```
END MODULE Constants
```

```
SUBROUTINE sorter (list)
```

```
    USE Constants
```

```
    CHARACTER(LEN=charlen), DIMENSION(:) :: list
```

```
    . . .
```

```
END SUBROUTINE sorter
```

```
USE Constants
```

```
CHARACTER(LEN=charlen) :: data(1000)
```

```
CALL sorter(data)
```

# Assumed Length CHARACTER

A **CHARACTER** length can be assumed  
The **length** is taken from the **actual argument**

You use an asterisk (\*) for the length  
It acts very like an **assumed shape array**

Note that it is a property of the **type**  
It is **independent** of any **array dimensions**

# Example (1)

```
FUNCTION is_palindrome (word)
  LOGICAL :: is_palindrome
  CHARACTER(LEN=*), INTENT(IN) :: word
  INTEGER :: N, I
  is_palindrome = .False.
  N = LEN(word)
comp: DO I = 1, (N-1)/2
      IF (word(I:I) /= word(N+1-I:N+1-I)) THEN
        RETURN
      END IF
    END DO comp
  is_palindrome = .True.
END FUNCTION is_palindrome
```



## Example (2)

Such **arguments** do not have to be **read-only**

```
SUBROUTINE reverse_word (word)
  CHARACTER(LEN=*), INTENT(INOUT) :: word
  CHARACTER(LEN=1) :: c
  N = LEN(word)
  DO I = 1, (N-1)/2
    c = word(I:I)
    word(I:I) = word(N+1-I:N+1-I)
    word(N+1-I:N+1-I) = c
  END DO
END SUBROUTINE reverse_word
```

# Character Workspace

The rules are very similar to those for arrays  
The length can be an almost arbitrary expression  
But it usually just shadows an argument

```
SUBROUTINE sort_words (words)
  CHARACTER(LEN=*) :: words(:)
  CHARACTER(LEN=LEN(words)) :: temp
  . . .
END SUBROUTINE sort_words
```

# Character Valued Functions

Functions can return **CHARACTER** values  
Fixed-length ones are the simplest

```
FUNCTION truth (value)
  IMPLICIT NONE
  CHARACTER(LEN=8) :: truth
  LOGICAL, INTENT(IN) :: value
  IF (value) THEN
    truth = '.True.'
  ELSE
    truth = '.False.'
  END IF
END FUNCTION truth
```

# Example

```
SUBROUTINE diagnose (message, value)
    CHARACTER(LEN=*), INTENT(IN) :: message
    REAL :: value
    PRINT *, message, value
END SUBROUTINE diagnose

CALL diagnose("Horrible failure",determinant)
```

# Static Data

Sometimes you need to store values locally  
Use a value in the next call of the procedure

- You do this with the **SAVE attribute**  
**Initialised variables** get that automatically  
It is good practice to specify it anyway

The best style avoids most such use  
It can cause trouble with **parallel** programming  
But it works, and lots of programs rely on it

# Example

This is a futile example, but shows the feature

```
SUBROUTINE Factorial (result)
  IMPLICIT NONE
  REAL, INTENT(OUT) :: result
  REAL, SAVE :: mult = 1.0, value = 1.0
  mult = mult+1.0
  value = value*mult
  result = value
END SUBROUTINE Factorial
```

# Warning

Omitting **SAVE** will usually **appear** to work  
But even a new compiler **version** may break it  
As will increasing the level of **optimisation**

- Decide which variables need it during **design**
- **Always** use **SAVE** if you want it  
And preferably never when you don't!
- **Never** assume it without specifying it

# Warning for C/C++ Users

Initialisation **without SAVE** initialises **once**  
It does **NOT** reinitialise **each time** it is called

- It can't be done using **Fortran initialisation**  
Do it using an **explicit assignment** statement



# Delayed Until Modules

Sometimes you need to share global data  
It's trivial, and can be done very cleanly

**Procedures** can be passed as **arguments**  
This is a very important facility for some people  
For historical reasons, this is a bit messy

- However, **internal procedures** can't be  
They can be in **Fortran 2008** – i.e. shortly

We will cover both of these under **modules**  
It just happens to be simplest that way!

# Other Features

There is a lot that we haven't covered  
We will return to some of it later

- The above covers the absolute basics  
Plus some other features you need to know
- Be a bit cautious when using other features  
Some have been omitted because of “**gotchas**”
- And I have over-simplified a few areas

# Extra Slides

Topics in the advanced slides on **procedures**

- **Argument association and updating**
- The semantics of **function calls**
- **Optional arguments**
- **Array- and character-valued functions**
- Mixing **explicit** and **assumed shape** arrays
- **Array arguments and sequence association**
- Miscellaneous other points

# Omissions

Rather a lot has been omitted here, unfortunately  
It's there in the notes, if you are interested

If you think that Fortran can't do it, look deeper  
Sorry about that, but this had to be simplified

# Introduction to Modern Fortran

*KIND, Precision and COMPLEX*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# The Basic Problem

**REAL** must be same size as **INTEGER**

This is for historical reasons – ask if you care

**32** bits allows **integers** of up to **2147483647**

Usually plenty for individual **array indices**

But floating-point **precision** is only **6** digits

And its **range** is only  **$10^{-38}$  –  $10^{+38}$**

Index values are not exact in floating-point

And there are many, serious numerical problems

# Example

```
REAL, DIMENSION(20000000) :: A  
REAL :: X  
X = SIZE(A)-1  
PRINT *, X
```

Prints 20000000.0 – which is not right  
That code needs only 80 MB to go wrong

See “How Computers Handle Numbers”  
Mainly on the numerical aspects

# Ordinary REAL Constants

These will often do what you expect

- But they will **very often** lose precision

0.0, 7.0, 0.25, 1.23, 1.23E12,  
0.1, 1.0E-1, 3.141592653589793

Only the first **three** will do what you expect

- In **old Fortran** constructs, can cause **chaos**  
E.g. as arguments to external libraries



# KIND Values

You can get the **KIND** of any expression

**KIND(var)** is the **KIND** value of **var**

**KIND(0.0)** is the **KIND** value of **REAL**

**KIND(0.0D0)** is that of **DOUBLE PRECISION**

This is described in a moment

**Implementation-dependent** integer values  
selecting the **type** (e.g. a specific **REAL**)

- Don't use **integer constants** directly

# SELECTED\_REAL\_KIND

You can request a minimum **precision** and **range**  
Both are specified in **decimal**

**SELECTED\_REAL\_KIND ( Prec [ , Range ] )**

This gives at least **Prec** decimal places  
and range  $10^{-Range} - 10^{+Range}$

E.g. **SELECTED\_REAL\_KIND(12)**  
at least **12** decimal places

# Warning: Time Warp

Unfortunately, we need to define a **module**  
We shall cover those quite a lot later

The one we shall define is trivial  
Just use it, and don't worry about the details  
Everything you need to know will be explained

Just compile it, but don't link it, using **-c**  
**nagfor -C=all -c double.f90**

# Using KIND (1)

You should write and compile a **module**

```
MODULE double
    INTEGER, PARAMETER :: DP =      &
        SELECTED_REAL_KIND(12)
END MODULE double
```

Immediately after every **procedure** statement

I.e. **PROGRAM**, **SUBROUTINE** or **FUNCTION**

```
USE double
IMPLICIT NONE
```

## Using KIND (2)

Declaring variables etc. is easy

```
REAL(KIND=DP) :: a, b, c  
REAL(KIND=DP), DIMENSION(10) :: x, y, z
```

Using constants is more tedious, but easy

```
0.0_DP, 7.0_DP, 0.25_DP, 1.23_DP, 1.23E12_DP,  
0.1_DP, 1.0E-1_DP, 3.141592653589793_DP
```

That's really all you need to know . . .

## Using KIND (3)

Note that the above makes it trivial to change  
**ALL** you need is to change the module

```
MODULE double
  INTEGER, PARAMETER :: DP = &
    SELECTED_REAL_KIND(15, 300)
END MODULE double
```

(15, 300) requires **IEEE 754 double** or better

Or even: `SELECTED_REAL_KIND(25, 1000)`

# DOUBLE PRECISION (1)

- The best way to control precision  
Most **flexible**, **portable** and **future-proof**  
Advisable if you **may** want to use **HECToR**

All older (**Fortran 77**) code will do it differently  
And quite a lot of programmers still do  
The old method is **fairly** reliable, today

- You need to know about this, but avoid it

# DOUBLE PRECISION (2)

DOUBLE PRECISION takes the space of 2 REALs  
⇒ It need not be any more accurate, though

- Almost always, REAL is 32-bit IEEE 754  
And DOUBLE PRECISION is 64-bit IEEE 754  
Precision is 15 digits, range is  $10^{-300} - 10^{+300}$

Main exception is Cray vector supercomputers  
And when using compiler options to change precision



# DOUBLE PRECISION (3)

You can use it just like **REAL** in declarations  
Using **KIND** is more modern and compact

```
REAL(KIND=KIND(0.0D0)) :: a, b, c
```

Constants use **D** for the exponent – **1.23D12** or **0.0D0**

```
REAL(KIND=KIND(0.0D0)) :: a, b, c
```

```
DOUBLE PRECISION, DIMENSION(10) :: x, y, z
```

```
0.0D0, 7.0D0, 0.25D0, 1.23D0, 1.23D12,  
0.1D0, 1.0D-1, 3.141592653589793D0
```

# Intrinsic Procedures

Almost all **intrinsic** 'just work' (i.e. are **generic**)  
**IMPLICIT NONE** removes most common traps

- Avoid specific (**old**) names for procedures  
**AMAX0**, **DMIN1**, **DSQRT**, **FLOAT**, **IFIX** etc.
- **DPROD** is also not **generic** – use a library
- Don't use the **INTRINSIC** statement
- Don't pass **intrinsic functions** as arguments

# Type Conversion (1)

This is the main “gotcha” – you should use

```
REAL(KIND=DP) :: x  
x = REAL(<integer expression>, KIND=DP)
```

Omitting the **KIND=DP** may lose precision

- With no warning from the compiler

**Automatic** conversion is actually safer!

```
x = <integer expression>  
x = SQRT(<integer expression>+0.0_DP)
```

# Type Conversion (2)

There is a **legacy** intrinsic function

If you are using explicit **DOUBLE PRECISION**

```
x = DBLE(<integer expression>)
```

All other “**gotchas**” are for **COMPLEX**

# Warning

You will often see code like:

```
REAL*8 X, Y, Z
```

```
INTEGER*8 M, N
```

- Most of the Web and many books are **wrong**

A **Fortran IV** feature, **NOT** a standard one  
'8' is **NOT** always the size in bytes

- I strongly recommend converting to **KIND**

# Old Fortran Libraries

Be **very** careful with external libraries

- Make sure **argument types** are right  
Automatic conversion does not happen  
Not will you get a diagnostic (in general)

Any **procedure** with no **explicit interface**  
I did say that using **old Fortran** was more painful

# INTEGER KIND

You can choose different sizes of integer

```
INTEGER, PARAMETER :: big = &  
    SELECTED_INT_KIND(12)  
INTEGER(KIND=big) :: bignum
```

**bignum** can hold values of up to at least  $10^{12}$   
Few users will need this – mainly for **OpenMP**

Some compilers may allocate smaller integers  
E.g. by using **SELECTED\_INT\_KIND(4)**

# CHARACTER KIND

It can be used to select the encoding

It is mainly a Fortran 2003 feature

Can select default, ASCII or ISO 10646

ISO 10646 is effectively Unicode

Useful for handling non-ASCII character sets

It is not covered in this course

Very few scientists want or use it



# Complex Arithmetic

Fortran is the answer – what was the question?

Has always been supported, and well integrated

**COMPLEX** is a (real, imaginary) pair of **REAL**

It uses the same **KIND** as underlying reals

```
COMPLEX(KIND=DP) :: c  
c = (1.23_DP, 4.56_DP)
```

Full range of operations, intrinsic functions etc.

# Example

```
COMPLEX(KIND=DP) :: c, d, e, f
```

```
c = (1.23_DP, 4.56_DP)*CONJG(d)+SIN(f*g)
```

```
e = EXP(d+c/f)*ABS(LOG(e))
```

The functions are the **complex** forms

E.g. **ABS** is  $\sqrt{re^2 + im^2}$

**CONJG** is **complex conjugate**, of course

Using **COMPLEX** really **IS** that simple!

# Worst “Gotcha”

- Must specify **KIND** in conversion function

```
c = CMPLX(<X-expr>, KIND=DP)
```

```
c = CMPLX(<X-expr>, <Y-expr>, KIND=DP)
```

This will not work – **KIND** is **default REAL**

Usually with no warning from the compiler

```
c = CMPLX(0.1_DP,0.2_DP)
```

# Conversion to REAL

```
REAL(KIND=DP) :: x  
COMPLEX(KIND=DP) :: c  
... lots of statements ...  
X = X+C  
c = 2.0_DP*x
```

Loses the imaginary part, without warning  
Almost all modern languages do the same

# A Warning for Old Code

```
C = DCMPLX(0.1_DP, 0.1_DP)
```

That is often seen in **Fortran IV** legacy code  
It doesn't work in **standard** (modern) Fortran

- It will be caught by **IMPLICIT NONE**

# Complex I/O

The form of I/O we have used is **list-directed**  
**COMPLEX** does what you would expect

```
COMPLEX(KIND=DP) :: c = (1.23_DP,4.56_DP)  
WRITE (*, *) C
```

Prints “(1.23,4.56)”  
And similarly for input

There is some more on **COMPLEX** I/O later

# Exceptions

Complex exceptions are **mathematically** hard

- **Overflow** often does what you won't expect  
Fortran, unfortunately, is no exception to this

See “**How Computers Handle Numbers**”

- Don't cause them in the first place
- Use the techniques described to detect them

# Introduction to Modern Fortran

## *Modules and Interfaces*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014



# Module Summary

- Similar to same term in other languages  
As usual, **modules** fulfil multiple purposes
- For shared declarations (i.e. “**headers**”)
- Defining **global data** (old **COMMON**)
- Defining **procedure interfaces**
- **Semantic extension** (described later)

And more ...

# Use Of Modules

- Think of a **module** as a **high-level interface**  
Collects **<whatevers>** into a coherent unit
- Design your **modules** carefully  
As the ultimate top-level **program structure**  
Perhaps only a few, perhaps dozens
- Good place for high-level comments  
**Please** document **purpose** and **interfaces**

# Module Interactions

Modules can **USE** other modules


Dependency graph shows **visibility**/usage

- **Modules** may not depend on themselves  
Languages that allow that are **very** confusing


Can do anything you are likely to get to work

- If you need to do more, ask for advice

# Module Dependencies



# Module Dependencies



# Module Structure

**MODULE** <name>

Static (often exported) data definitions

**CONTAINS**

Procedure definitions (i.e. their code)

**END MODULE** <name>

**Files** may contain several **modules**

**Modules** may be split across many **files**

- For simplest use, keep them **1≡1**

# IMPLICIT NONE

Add **MODULE** to the places where you use this

```
MODULE double
  IMPLICIT NONE
  INTEGER, PARAMETER :: DP = KIND(0.0D0)
END MODULE double
```

```
MODULE parameters
  USE double
  IMPLICIT NONE
  REAL(KIND=DP), PARAMETER :: one = 1.0_DP
END MODULE parameters
```

# Reminder

I do not always do it, because of space



# Example (1)

```
MODULE double
    INTEGER, PARAMETER :: DP = KIND(0.0D0)
END MODULE double
```

```
MODULE parameters
    USE double
    REAL(KIND=DP), PARAMETER :: one = 1.0_DP
    INTEGER, PARAMETER :: NX = 10, NY = 20
END MODULE parameters
```

```
MODULE workspace
    USE double ;    USE parameters
    REAL(KIND=DP), DIMENSION(NX, NY) :: now, then
END MODULE workspace
```

## Example (2)

The **main program** might use them like this

```
PROGRAM main
  USE double
  USE parameters
  USE workspace
  . . .
END PROGRAM main
```

- Could omit the **USE double** and **USE parameters**  
They would be **inherited** through **USE workspace**

# Shared Constants

We have already seen and used this:

```
MODULE double  
    INTEGER, PARAMETER :: DP = KIND(0.0D0)  
END MODULE double
```

You can do a great deal of that sort of thing

- Greatly improves **clarity** and **maintainability**  
The larger the program, the more it helps

# Example

```
MODULE hotchpotch
  INTEGER, PARAMETER :: DP = KIND(0.0D0)
  REAL(KIND=DP), PARAMETER ::  &
    pi = 3.141592653589793_DP,  &
    e = 2.718281828459045_DP
  CHARACTER(LEN=*), PARAMETER ::  &
    messages(3) =  &
      (\ "Hello", "Goodbye", "Oh, no!" \)
  INTEGER, PARAMETER :: stdin = 5, stdout = 6
  REAL(KIND=DP), PARAMETER,  &
    DIMENSION(0:100, -1:25, 1:4) :: table =  &
      RESHAPE( (/ . . . /), (/ 101, 27, 4 /) )
END MODULE hotchpotch
```

# Global Data

Variables in modules define global data  
These can be fixed-size or allocatable arrays

- You need to specify the **SAVE** attribute  
Set automatically for initialised variables  
But it is good practice to do it explicitly

A simple **SAVE** statement saves everything

- That isn't always the best thing to do

# Example (1)

```
MODULE state_variables
  INTEGER, PARAMETER :: nx=100, ny=100
  REAL, DIMENSION(NX, NY), SAVE ::  &
    current, increment, values
  REAL, SAVE :: time = 0.0
END MODULE state_variables

USE state_variables
IMPLICIT NONE
DO
  current = current + increment
  CALL next_step(current, values)
END DO
```

## Example (2)

This is equivalent to the previous example

```
MODULE state_variables
  IMPLICIT NONE
  SAVE
  INTEGER, PARAMETER :: nx=100, ny=100
  REAL, DIMENSION(NX, NY) :: &
    current, increment, values
  REAL :: time = 0.0
END MODULE state_variables
```

## Example (3)

The sizes do not have to be fixed

```
MODULE state_variables
    REAL, DIMENSION(:, :), ALLOCATABLE, &
        SAVE :: current, increment, values
END MODULE state_variables
```

```
USE state_variables
IMPLICIT NONE
INTEGER :: NX, NY
READ *, NX, NY
ALLOCATE (current(NX, NY), increment(NX, NY), &
    values(NX, NY))
```



# Use of SAVE

If a **variable** is set in one **procedure**  
and then it is used in another

- You must specify the **SAVE** attribute

- If not, **very** strange things **may** happen

If will usually “**work**”, under most compilers

A new version will appear, and then it won't

- Applies if the **association** is via the **module**

Not when it is passed as an **argument**

# Example (1)

```
MODULE status  
    REAL :: state  
END MODULE status
```

```
SUBROUTINE joe  
    USE status  
    state = 0.0  
END SUBROUTINE joe
```

```
SUBROUTINE alf (arg)  
    REAL :: arg  
    arg = 0.0  
END SUBROUTINE alf
```

## Example (2)

```
SUBROUTINE fred
  USE status

  CALL joe
  PRINT *, state    ! this is UNDEFINED

  CALL alf(state)
  PRINT *, state    ! this is defined to be 0.0

END SUBROUTINE fred
```

# Shared Workspace

Shared scratch space can be useful for HPC  
It can avoid excessive memory fragmentation

You can omit **SAVE** for simple scratch space  
This can be significantly more efficient

- Design your data use carefully  
Separate global scratch space from storage  
And use them consistently and correctly
- This is good practice in any case

# Module Procedures (1)

Procedures now need explicit interfaces

E.g. for assumed shape or keywords

Without them, must use Fortran 77 interfaces

- Modules are the primary way of doing this

We will come to the secondary one later

Simplest to include the procedures in modules

The procedure code goes after CONTAINS

This is what we described earlier

# Example

```
MODULE mymod
CONTAINS
    FUNCTION Variance (Array)
        REAL :: Variance, X
        REAL, INTENT(IN), DIMENSION(:) :: Array
        X = SUM(Array)/SIZE(Array)
        Variance = SUM((Array-X)**2)/SIZE(Array)
    END FUNCTION Variance
END MODULE mymod

PROGRAM main
    USE mymod
    ...
    PRINT *, 'Variance = ', Variance(array)
```

# Module Procedures (2)

- **Modules** can contain any number of **procedures**
- You can use any number of **modules**

```
PROGRAM main
  USE mymod
  REAL, DIMENSION(10) :: array
  PRINT *, 'Type 10 values'
  READ *, array
  PRINT *, 'Variance = ', Variance(array)
END PROGRAM main
```

# Using Procedures

Internal procedures or module procedures?

Use either technique for solving test problems

- They are the best techniques for real code  
Simplest, and give full access to functionality  
We will cover some other ones later

- Note that, if a procedure is in a module  
it may still have internal procedures



# Example

```
MODULE mymod
CONTAINS
    SUBROUTINE Sorter (array, opts)
        . . .
    CONTAINS
        FUNCTION Compare (value1, value2, flags)
            . . .
        END FUNCTION Compare
        SUBROUTINE Swap (loc1, loc2)
            . . .
        END FUNCTION Swap
    END SUBROUTINE Sorter
END MODULE mymod
```

# Procedures in Modules (1)

That is including all **procedures** in **modules**  
Works very well in almost all programs

- There really isn't much more to it

It doesn't handle very large modules well  
Try to avoid designing those, if possible

It also doesn't handle **procedure arguments**  
Unfortunately, doing that has had to be omitted

# Procedures in Modules (2)

They are very like **internal procedures**

Everything accessible in the **module**  
can also be used in the **procedure**

Again, a **local name** takes precedence  
But reusing the same name is very confusing

# Procedures in Modules (3)

```
MODULE thing
    INTEGER, PARAMETER :: temp = 123
CONTAINS
    SUBROUTINE pete ()
        INTEGER, PARAMETER :: temp = 456
        PRINT *, temp
    END SUBROUTINE pete
END MODULE thing
```

Will print 456, not 123

Avoid doing this – it's very confusing

# Derived Type Definitions

We shall cover these later:

```
MODULE Bicycle
  TYPE Wheel
    INTEGER :: spokes
    REAL    :: diameter, width
    CHARACTER(LEN=15) :: material
  END TYPE Wheel
END MODULE Bicycle

USE Bicycle
TYPE(Wheel) :: w1
```

# Compiling Modules (1)

This is a **FAQ** – Frequently Asked Question  
The problem is the **answer** isn't simple

- That is why I give some of the advice that I do

The following advice will **not** always work  
OK for **most** compilers, but **not** necessarily **all**

- This is **only** the **Fortran module** information

## Compiling Modules (2)

The **module name** need not be the **file name**  
Doing that is strongly recommended, though

- You can include any number of **whatevers**

You now **compile** it, but don't **link** it  
**nagfor -C=all -c mymod.f90**

It will create files like **mymod.mod** and **mymod.o**  
They contain the **interface** and the **code**

Will describe the process in more detail later

# Using Compiled Modules

All the program needs is the **USE** statements

- Compile all of the modules in a **dependency order**  
If **A** contains **USE B**, compile **B** first
- Then add a **\*.o** for every module when **linking**

```
nagfor -C=all -o main main.f90 mymod.o
```

```
nagfor -C=all -o main main.f90 \  
    mod_a.o mod_b.o mod_c.o
```



# Take a Breather

That is most of the basics of **modules**  
Except for **interfaces** and **access control**

The **first** question covers the material **so far**

The remainder is **important** and **useful**  
But it is unfortunately rather more **complicated**

# What Are Interfaces?

The **FUNCTION** or **SUBROUTINE** statement

And everything **directly connected** to that

**USE** if needed for **argument declarations**

- And don't forget a **function result** declaration

Strictly, the **argument names** are not part of it

You are **strongly** advised to keep them the same

Which **keywords** if the **interface** and **code** differ?

Actually, it's the ones in the **interface**

# Interface Blocks

These start with an **INTERFACE** statement  
Include any number of **procedure interfaces**  
And end with an **END INTERFACE** statement

```
INTERFACE
  SUBROUTINE Fred (arg)
    REAL :: arg
  END FUNCTION Fred
  FUNCTION Joe ()
    LOGICAL :: Joe
  END FUNCTION Joe
END INTERFACE
```

# Example

```
SUBROUTINE CHOLESKY (A)  ! this is part of it
  USE errors             ! this ISN'T part of it
  USE double             ! this is, because of A
  IMPLICIT NONE          ! this ISN'T part of it
  INTEGER :: J, N        ! this ISN'T part of it
  REAL(KIND=dp) :: A(:, :), X  ! A is but not X
```

...

```
END SUBROUTINE CHOLESKY
```

```
INTERFACE
```

```
  SUBROUTINE CHOLESKY (A)
```

```
    USE double
```

```
    REAL(KIND=dp) :: A(:, :)
```

```
  END SUBROUTINE CHOLESKY
```

```
END INTERFACE
```

# Interfaces In Procedures

Can use an **interface block** as a **declaration**  
Provides an **explicit interface** for a **procedure**

Can be used for ordinary procedure calls  
But using **modules** is almost always better

- It is essential for **procedure arguments**  
Can't put a **dummy argument name** in a **module**!

More on this in the **Make and Linking** lecture

# Example (1)

Assume this is in **module application**

```
FUNCTION apply (arr, func)
  REAL :: apply, arr(:)
  INTERFACE
    FUNCTION func (val)
      REAL :: func, val
    END FUNCTION
  END INTERFACE
  apply = 0.0
  DO I = 1,UBOUND(arr, 1)
    apply = apply + func(val = arr(i))
  END DO
END FUNCTION apply
```

## Example (2)

And these are in **module functions**

```
FUNCTION square (arg)
  REAL :: square, arg
  square = arg**2
END FUNCTION square
```

```
FUNCTION cube (arg)
  REAL :: cube, arg
  cube = arg**3
END FUNCTION cube
```

## Example (3)

```
PROGRAM main
  USE application
  USE functions
  REAL, DIMENSION(5) :: A = (/ 1.0, 2.0, 3.0, 4.0, 5.0 /)
  PRINT *, apply(A,square)
  PRINT *, apply(A,cube)
END PROGRAM main
```

Will produce something like:

```
55.0000000
2.2500000E+02
```



# Interface Bodies and Names (1)

An **interface body** does **not** import names  
The reason is that you can't **undeclare** names

For example, this does not work as expected:

```
USE double      ! This doesn't help
INTERFACE
    FUNCTION square (arg)
        REAL(KIND=dp) :: square, arg
    END FUNCTION square
END INTERFACE
```

# Interface Bodies and Names (2)

So there is **another statement** to import names:

```
USE double
INTERFACE
    FUNCTION square (arg)
        IMPORT :: dp      ! This solves it
        REAL(KIND=dp) :: square, arg
    END FUNCTION square
END INTERFACE
```

It is available **only** in **interface bodies**

# Accessibility (1)

Can separate **exported** from **hidden** definitions

Fairly easy to use in simple cases

- Worth considering when designing modules

**PRIVATE** **names** accessible only in **module**

I.e. in **module procedures** after **CONTAINS**

**PUBLIC** **names** are accessible by **USE**

This is commonly called **exporting** them

# Accessibility (2)

They are just another **attribute** of declarations

```
MODULE fred
  REAL, PRIVATE :: array(100)
  REAL, PUBLIC :: total
  INTEGER, PRIVATE :: error_count
  CHARACTER(LEN=50), PUBLIC :: excuse
CONTAINS
  . . .
END MODULE fred
```

# Accessibility (3)

**PUBLIC/PRIVATE** statement sets the default  
The default default is **PUBLIC**

```
MODULE fred
  PRIVATE
  REAL :: array(100)
  REAL, PUBLIC :: total
CONTAINS
  . . .
END MODULE fred
```

Only **TOTAL** is accessible by **USE**

# Accessibility (4)

You can specify **names** in the **statement**  
Especially useful for **included names**

```
MODULE workspace  
  USE double  
  PRIVATE :: DP  
  REAL(KIND=DP), DIMENSION(1000) :: scratch  
END MODULE workspace
```

**DP** is no longer **exported** via **workspace**

# Partial Inclusion (1)

You can include only some **names** in **USE**

**USE** bigmodule, **ONLY** : errors, invert

Makes only **errors** and **invert** visible

However many **names** **bigmodule** exports

Using **ONLY** is good practice

Makes it easier to keep track of uses

Can find out what is used where with **grep**

# Partial Inclusion (2)

- One case when it is **strongly** recommended  
When using **USE** in **modules**
- All **included names** are **exported**  
Unless you explicitly mark them **PRIVATE**
- Ideally, use both **ONLY** and **PRIVATE**  
Almost always, use **at least one** of them
- Another case when it is **almost essential**  
Is if you don't use **IMPLICIT NONE** religiously



# Partial Inclusion (3)

If you don't restrict **exporting** and **importing**:

A typing error could trash a **module variable**

Or forget that you had already used the **name**  
In another **file** far, far away ...

- The resulting chaos is almost unfindable  
From bitter experience – in Fortran and **C**!

# Example (1)

MODULE settings

INTEGER, PARAMETER :: DP = KIND(0.0D0)

REAL(KIND=DP) :: Z = 1.0\_DP

END MODULE settings

MODULE workspace

USE settings

REAL(KIND=DP), DIMENSION(1000) :: scratch

END MODULE workspace

## Example (2)

```
PROGRAM main
  IMPLICIT NONE
  USE workspace
  Z = 123
  . . .
END PROGRAM main
```

- DP is inherited, which is OK
- Did you mean to update Z in settings?

No problem if workspace had used ONLY : DP

## Example (3)

The following are **better** and **best**

```
MODULE workspace
```

```
    USE settings, ONLY : DP
```

```
    REAL(KIND=DP), DIMENSION(1000) :: scratch
```

```
END MODULE workspace
```

```
MODULE workspace
```

```
    USE settings, ONLY : DP
```

```
    PRIVATE :: DP
```

```
    REAL(KIND=DP), DIMENSION(1000) :: scratch
```

```
END MODULE workspace
```

# Renaming Inclusion (1)

You can rename a **name** when you **include** it

**WARNING:** this is footgun territory  
[ i.e. point gun at foot; pull trigger ]

This technique is sometimes **incredibly** useful

- But is always **incredibly dangerous**

Use it only when you **really** need to  
And even then **as little** as possible

## Renaming Inclusion (2)

```
MODULE corner  
    REAL, DIMENSION(100) :: pooh  
END MODULE corner
```

```
PROGRAM house  
    USE corner, sanders => pooh  
    INTEGER, DIMENSION(20) :: pooh  
    ...  
END PROGRAM house
```

**pooh** is accessible under the **name sanders**  
The **name pooh** is the **local array**

# Why Is This Lethal?

```
MODULE one  
    REAL :: X  
END MODULE one
```

```
MODULE two  
    USE one, Y => X  
    REAL :: Z  
END MODULE two
```

```
PROGRAM three  
    USE one ;    USE two  
    ! Both X and Y refer to the same variable  
END PROGRAM three
```

# Interfaces and Access Control

These are things that have been **omitted**

They're there in the notes, if you are interested

They are extremely important for large programs

But time is too tight to teach them now

- Do only the first practical and skip the rest



# Introduction to Modern Fortran

## *Derived Types*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Summary

There is one important **new feature** to cover

It is **not complicated**, as we shall do it

- But we won't cover it in great depth

Doing it **fully** would be a course in itself

The same applies in **other** languages, too

# What Are Derived Types?

As usual, a **hybrid** of two, unrelated concepts **C++**, **Python** etc. are very similar

- One is **structures** – i.e. composite objects  
Arbitrary **types**, statically indexed by name

- The other is **user-defined types**  
Often called **semantic extension**  
This is where **object orientation** comes in

- This course will describe only the **former**

# Why Am I Wimping Out?

Fortran 2003 has really changed this  
full object orientation  
semantic extension  
polymorphism (abstract types)  
and lots more

The course was already getting too big  
And, yes, I was getting sick of writing it!

This area justifies a separate course  
About one day or two afternoons, not three days  
Please ask if you would like it written

# Simple Derived Types

```
TYPE :: Wheel
    INTEGER :: spokes
    REAL    :: diameter, width
    CHARACTER(LEN=15) :: material
END TYPE Wheel
```

That defines a **derived type** **Wheel**  
Using **derived types** needs a special syntax

```
TYPE(Wheel) :: w1
```

# More Complicated Ones

You can include almost anything in there

```
TYPE :: Bicycle
  CHARACTER(LEN=80) :: description(100)
  TYPE(Wheel) :: front, back
  REAL, ALLOCATABLE, DIMENSION(:) :: times
  INTEGER, DIMENSION(100) :: codes
END TYPE Bicycle
```

And so on ...

# Fortran 95 Restriction

Fortran 95 was much more restrictive

You couldn't have **ALLOCATABLE** arrays

You had to use **POINTER** instead

Fortran 2003 removed that restriction

You may come across **POINTER** in old code

It can usually be replaced by **ALLOCATABLE**

Ask if you hit problems and want to check

# Component Selection

The selector ‘%’ is used for this  
Followed by a **component** of the **derived type**

It delivers whatever **type** that **field** is  
You can then **subscript** or **select** it

```
TYPE(Bicycle) :: mine
```

```
mine%times(52:53) = (/ 123.4, 98.7 /)  
PRINT *, mine%front%spokes
```



# Selecting from Arrays

You can **select** from **arrays** and **array sections**  
It produces an **array** of that **component** alone

```
TYPE :: Rabbit
    CHARACTER(LEN=16) :: variety
    REAL :: weight, length
    INTEGER :: age
END TYPE Rabbit
TYPE(Rabbit), DIMENSION(100) :: exhibits
REAL, DIMENSION(50) :: fattest

fattest = exhibits(51:)%weight
```

# Assignment (1)

You can **assign** complete **derived types**  
That copies the value element-by-element

```
TYPE(Bicycle) :: mine, yours
```

```
yours = mine
```

```
mine%front = yours%back
```

**Assignment** is the only **intrinsic operation**

You can redefine that or define other operations  
But they are some of the topics I am omitting

# Assignment (2)

Each **derived type** is a separate type  
You **cannot** assign between different ones

```
TYPE :: Fred
    REAL :: x
END TYPE Fred
TYPE :: Joe
    REAL :: x
END TYPE Joe
TYPE(Fred) :: a
TYPE(Joe) :: b
a = b    ! This is erroneous
```

# Constructors

A constructor creates a derived type value

```
TYPE Circle  
    REAL :: X, Y, radius  
    LOGICAL :: filled  
END TYPE Circle
```

```
TYPE(Circle) :: a  
a = Circle(1.23, 4.56, 2.0, .False.)
```

Or use keywords for components (Fortran 2003)

```
a = Circle(X = 1.23, Y = 4.56, radius = 2.0, filled = .False.)
```

# Default Initialisation

You can specify default **initial values**

```
TYPE :: Circle  
    REAL :: X = 0.0, Y = 0.0, radius = 1.0  
    LOGICAL :: filled = .False.  
END TYPE Circle
```

```
TYPE(Circle) :: a, b, c  
a = Circle(1.23, 4.56, 2.0, .True.)
```

This becomes much more useful with **keywords**

```
a = Circle(X = 1.23, Y = 4.56)
```

# I/O on Derived Types

Can do normal I/O with the **ultimate components**

A **derived type** is flattened much like an array  
[ recursively, if it includes **derived types** ]

```
TYPE(Circle) :: a, b, c
```

```
a = Circle(1.23, 4.56, 2.0, .True.)
```

```
PRINT *, a ; PRINT *, b ; PRINT *, c
```

```
1.2300000  4.5599999  2.0000000  T
```

```
0.0000000E+00  0.0000000E+00  1.0000000  F
```

```
0.0000000E+00  0.0000000E+00  1.0000000  F
```

# Private Derived Types

When you define them in **modules**

A **derived type** can be **wholly private**  
I.e. accessible only to **module procedures**

Or its **components** can be **hidden**  
I.e. it's visible as an **opaque type**

Both useful, even without **semantic extension**

# Wholly Private Types

```
MODULE Marsupial
  TYPE, PRIVATE :: Wombat
    REAL :: weight, length
  END TYPE Wombat
  REAL, PRIVATE :: Koala
CONTAINS
  . . .
END MODULE Marsupial
```

**Wombat** is not **exported** from **Marsupial**  
No more than the **variable** **Koala** is



# Hidden Components (1)

```
MODULE Marsupial
  TYPE :: Wombat
    PRIVATE
    REAL :: weight, length
  END TYPE Wombat
CONTAINS
  ...
END MODULE Marsupial
```

Wombat **IS** exported from Marsupial  
But its **components** (**weight**, **length**) are not

# Hidden Components (2)

Hidden components allow opaque types

The module procedures use them normally

- Users of the module can't look inside them

They can assign them like variables

They can pass them as arguments

Or call the module procedures to work on them

An important software engineering technique

Usually called data encapsulation

# Trees

E.g. type **A** contains an array of type **B**  
Objects of type **B** contain arrays of type **C**

```
TYPE :: Leaf
    CHARACTER(LEN=20) :: name
    REAL(KIND=dp), DIMENSION(3) :: data
END TYPE Leaf
TYPE :: Branch
    TYPE(Leaf), ALLOCATABLE :: leaves(:)
END TYPE Branch
TYPE :: Trunk
    TYPE(Branch), ALLOCATABLE :: branches(:)
END TYPE Trunk
```

# Recursive Types

**Pointers** allow that to be done a little more flexibly  
You don't need a separate type for **each level**

People often use more complicated structures  
You build those using **derived types**  
E.g. **linked lists** (also called **chains**)

Both very commonly used for **sparse matrices**  
And algorithms like **Dirichlet tessellation**

We shall return to this when we cover **pointers**

# Opaque Types etc.

This is another using aspect that has been omitted  
It's there in the notes, if you are interested

- Skip the practical that needs that facility

# Introduction to Modern Fortran

## *I/O and Files*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# I/O Generally

Most descriptions of I/O are only **half-truths**  
Those work most of the time – until they blow up  
Most modern language **standards** are like that

Fortran is rather better, but there are downsides  
**Complexity** and **restrictions** being two of them

- Fortran is **much easier to use** than it **seems**
- This is about what you can rely on **in practice**

We will start with the basic principles

# Some ‘Recent’ History

Fortran I/O (**1950s**) predates even **mainframes**  
**OPEN** and **filenames** was a **CC†** of **c. 1975**

**Unix/C** spread through CS depts **1975–1985**

**ISO C**’s I/O model was a **CC†** of **1985–1988**

Modern languages use the **C/POSIX** I/O model

Even **Microsoft** systems are like **Unix** here

- The I/O models have little in common

† **CC** = **committee compromise**



# Important Warning

It is often **better** than **C/C++** and often **worse**  
But it is **very** different at **all levels**

- It is **critical** not to think in **C**-like terms  
Trivial **C/C++** tasks may be infeasible in Fortran

As always, use the **simplest** code that works  
Few people have much trouble if they do that

- Ask for help with any problems here

# Fortran's Sequential I/O Model

A **Unix** file is a sequence of **characters** (bytes)

A **Fortran** file is a sequence of **records** (lines)

For simple, text use, these are almost equivalent

In both **Fortran** and **C/Unix**:

- Keep text records short (say, **< 250 chars**)
- Use only **printing characters** and **space**
- Terminate all lines with a plain **newline**
- **Trailing spaces** can appear and disappear

# What We Have Used So Far

To remind you what you have been doing so far:

`PRINT *`, can be written `WRITE (*,*)`

`READ *`, can be written `READ (*,*)`

`READ/WRITE (*,*)` is shorthand for  
`READ/WRITE (UNIT=*, FMT=*)`

`READ *`, ... and `PRINT *`, ... are **legacies**

Their syntax is **historical** and **exceptional**

# Record-based I/O

- Each **READ** and **WRITE** uses **1+** records  
Any unread characters are skipped for **READ**  
**WRITE** ends by writing an **end-of-line** indicator
  - Think in terms of units of whole lines  
A **WRITE** builds one or more **whole** lines  
A **READ** consumes one or more **whole** lines
- Fortran 2003** relaxes this, to **some** extent

# Fortran's I/O Primitives

All Fortran I/O is done with special **statements**

Any **I/O procedure** is a **compiler extension**

Except as above, all of these have the syntax:

**<statement> (<control list>) <transfer list>**

The **<transfer list>** is only for **READ** and **WRITE**

The **<control list>** items have the syntax:

**<specifier>=<value>**

# Translation

An I/O statement is rather like a **command**

A **<control list>** is rather like a set of **options**  
Though not all of the **specifiers** are optional

The **<transfer list>** is a list of **variables** to read  
or a list of **expressions** to write

We now need to describe them in more detail

# Specifier Values (1)

All **specifier values** can be **expressions**

If they return values, they must be **variables**

- Except for **\*** in **UNIT=\*** or **FMT=\***

Even lunatic code like this is permitted

```
INTEGER, DIMENSION(20) :: N  
CHARACTER(LEN=50) :: C
```

```
WRITE (UNIT = (123*K)/56+2, FMT = C(3:7)//'),    &  
      IOSTAT=N(J**5-15))
```

## Specifier Values (2)

The examples will usually use explicit constants

```
OPEN (23, FILE='trace.out', RECL=250)
```

- But you are advised to **parameterise** units  
And anything else that is **system dependent**  
Or you might need to change later

```
INTEGER, PARAMETER :: tracing = 23, tracelen = 250  
CHARACTER(LEN=*), PARAMETER :: &  
    tracefile = 'trace.out'
```

```
OPEN (tracing, FILE=tracefile, RECL=tracelen)
```



# Basics of READ and WRITE

READ/WRITE (<control list>) <transfer list>

Control items have form <specifier> = <value>

UNIT is the only compulsory control item

The UNIT= can be omitted if the unit comes first

The unit is an integer identifying the connection

It can be an expression, and its value is used

UNIT=\* is an exceptional syntax

It usually means stdin and stdout

# Transfer Lists

A **list** is a comma-separated sequence of items  
The list may be empty (**READ and WRITE**)

- A basic **output** item is an **expression**
- A basic **input** item is a **variable**

**Arrays** and **array expressions** are allowed

- They are expanded in **array element order**

Fancy **expressions** will often cause a copy  
**Array sections** **should not** cause a copy

# Example

```
REAL :: X(10)  
READ *, X(:7)  
PRINT *, X(4:9:3)*1000.0
```

1.23 2.34 3.45 4.56 5.67 6.78 7.89

Produces a result like:

4.5600000E+03 7.8900000E+03

# Empty Transfer Lists

These are allowed, defined and meaningful

`READ (*, *)` skips the next line

`WRITE (*, *)` prints a blank line

`WRITE (*, FORMAT)` prints any text in `FORMAT`

That may print **several lines**

# A Useful Trick

A useful and fairly common construction

```
INTEGER :: NA  
REAL, DIMENSION(1:100) :: A  
READ *, NA, A(1:NA)
```

Fortran evaluates a transfer list as it executes it

- **Be warned:** easy to exceed array bounds

At least, you should check the length afterwards  
Safer to put on separate lines and check first

# Implied DO-loops

There is an **alternative** to array expressions  
Equivalent, but older and often more convenient

Items may be ( **<list>** , **<indexed loop control>** )  
This repeats in the **loop** order (just like **DO**)

( ( A(I,J) , J = 1,3 ) , B(I), I = 6,2,-2 )

A(6,1), A(6,2), A(6,3), B(6), A(4,1), A(4,2),  
A(4,3), B(4), A(2,1), A(2,2), A(2,3), B(2)

# Programming Notes

You can do I/O of **arrays** in three ways:

- You can write a **DO**-loop around the I/O
- **Array expressions** for selecting and ordering
- You can use **implied DO-loops**

Use whichever is most **convenient** and **clearest**

There are no problems with combining them

More examples of their use will be shown later

There isn't a general ranking of efficiency

# The UNIT Specifier

- A **unit** is an **integer value**  
Except for **UNIT=\***, described above  
It identifies the **connection** to a **file**
- **UNIT=** can be omitted if the **unit** is **first**

A **unit** must be **connected** to a **file** before use  
Generally use values in the range **10–99**

- That's all you need to know for now



# The FMT Specifier

This sets the type of I/O and must match the file

- **FMT=** can be omitted if the **format** is **second** and the **first** item is the **unit**
- **FMT=\*** indicates **list-directed** I/O
- **FMT=<format>** indicates **formatted** I/O

These can be interleaved on **formatted** files

- No **FMT** specifier indicates **unformatted** I/O

# Example

These are **formatted** I/O statements

```
WRITE (UNIT = *, FMT = '(2F5.2)') c
```

```
READ (99, '(F5.0)') x
```

```
WRITE (*, FMT = myformat) p, q, r
```

These are **list-directed** I/O statements

```
WRITE (UNIT = *, FMT = *) c
```

```
READ (99, *) x
```

These are **unformatted** I/O statements

```
WRITE (UNIT = 64) c
```

```
READ (99) x
```

# List-Directed Output (1)

What you have been doing with 'PRINT \*,'

The **transfer list** is split into basic elements  
Each element is then formatted appropriately  
It is separated by **spaces** **and/or** a **comma**

- Except for **adjacent CHARACTER** items  
Write spaces **explicitly** if you want them

The **format** and **layout** are compiler-dependent

# Example

```
REAL :: z(3) = (/4.56, 4.56, 4.56/)
CHARACTER(LEN=1) :: c = 'a'
PRINT *, 1.23, 'Oh dear', z, c, '"', c, ' ', c, c
```

Produces (under one compiler):

```
1.2300000 Oh dear 4.5599999 4.5599999
4.5599999 a"a aa
```

## List-Directed Output (2)

You can cause **character strings** to be quoted  
Very useful if **writing** data for **reinput**

```
WRITE (11, *, DELIM='quote') 'Kilroy was here'
```

```
"Kilroy was here"
```

Also **DELIM='apostrophe'** and **DELIM='none'**

They can also be specified in **OPEN**

Apply to all **WRITE**s with no **DELIM**

# List-Directed Input (1)

What you have been doing with 'READ \*,'

This does the reverse of 'PRINT \*,'

The closest Fortran comes to free-format input

- It automatically checks the data type
  - OK for lists of numbers and similar
- Not much good for genuinely free-format

# List-Directed Input (2)

Strings may be **quoted**, or not

Using either **quote** (") or **apostrophe** (')

- Quote **all** strings containing the following:  
    ,    /    "    '    \*    **space**    **end-of-line**

For the reasons why, read the specification

**List-directed** input is actually quite powerful

But very unlike all other modern languages

# Example

```
REAL :: a, b, c  
CHARACTER(LEN=8) :: p, q  
READ *, a, p, b, q, c  
PRINT *, a, p, b, q, c
```

```
123e-2 abcdefghijkl -003 "P""Q'R" 4.56
```

Produces (under one compiler):

```
1.23000000 abcdefgh -3.00000000 P"Q'R"  
4.55999999
```



# Free-Format

Free-format I/O is not traditional in Fortran

Formatted output is far more flexible

Fortran 2003 adds some free-format support

Free-format input can be very tricky in Fortran

But it isn't hard to read lists of numbers

There is some more on this in extra slides

# Unformatted I/O is Simple

Very few users have any trouble with it

- It is NOT like C binary I/O
- It is unlike anything in C

Most problems come from “thinking in C”

# Unformatted I/O (1)

- It is what you use for saving data in files  
E.g. writing your own **checkpoint/restart**  
Or transferring bulk data between programs
- No formatting/decoding makes it a **lot** faster  
**100+** times less CPU time has been observed
- Assume **same hardware** and **same system**  
If not, see other courses and **ask for help**

# Unformatted I/O (2)

Just reads and writes data as stored in memory

- You must read back into the **same types**
- Each transfer uses **exactly one record**

With extra **control data** for record boundaries

You don't need to know what it looks like

- Specify **FORM='unformatted'** in **OPEN**  
**stdin**, **stdout** and **terminals** are **not** suitable

That's **ALL** that you absolutely need to know!

# Example

```
INTEGER, DIMENSION(1000) :: index  
REAL, DIMENSION(1000000) :: array
```

```
OPEN (31, FILE='fred', FORM='unformatted')
```

```
DO k = 1,...  
    WRITE (31) k, m, n, index(:m), array(:n)  
END DO
```

In another run of the program, or after rewinding:

```
DO k = 1,...  
    READ (31) junk, m, n, index(:m), array(:n)  
END DO
```

# Programming Notes

- Make each **record** (i.e. **transfer**) quite large  
But don't go over **2 GB** per record

- I/O with whole arrays is generally fastest

```
INTEGER :: N(1000000)  
READ (29) N
```

Array sections **should** be comparably fast

- Remember about checking for **copying**
- **Implied DO-loops** should be avoided  
At least for large **loop counts**

# Formatted I/O

**READ** or **WRITE** with an explicit **format**

A **format** is just a **character string**

It can be specified in any one of three ways:

- A **CHARACTER** expression
- A **CHARACTER** array  
Concatenated in **array element order**
- The **label** of a **FORMAT** statement  
Old-fashioned, and best avoided

# Formats (1)

A **format** is **items** inside **parentheses**  
**Blanks** are ignored, except in **strings**

`' ( i3,f 5 . 2) '`  $\equiv$  `'(i3,f5.2)'`

We will see why this is so useful later

Almost any **item** may have a **repeat count**

`'(3 i3, 2 f5.2)'`  $\equiv$  `'(i3, i3, i3, f5.2, f5.2)'`



## Formats (2)

A **group** of **items** is itself an **item**

**Groups** are enclosed in **parentheses**

E.g. ‘( 3 (2 i3, f5.2 ) )’ expands into:

‘(i3, i3, f5.2, i3, i3, f5.2, i3, i3, f5.2)’

Often used with **arrays** and **implied DO-loops**

Nesting them deeply can be confusing

# Example

```
REAL, DIMENSION(2, 3) :: coords  
INTEGER, DIMENSION(3) :: index
```

```
WRITE (29, '( 3 ( i3, 2 f5.2 ) )') &  
    ( index(i), coords(:, i), i = 1,3)
```

This is how to use a **CHARACTER** constant:

```
CHARACTER(LEN=*), PARAMETER :: &  
    format = '( 3 ( i3, 2 f5.2 ) )'
```

```
WRITE (29, format) ( index(i), coords(:, i), i = 1,3)
```

# Transfer Lists And Formats

Logically, both are expanded into flat lists  
I.e. sequences of basic items and descriptors

The transfer list is the primary one  
Basic items are taken from it one by one  
Each then matches the next edit descriptor

The item and descriptor must be compatible  
E.g. REAL vars must match REAL descs

# Input Versus Output

We shall mainly describe formatted **output**  
This is rather simpler and more general

Unless mentioned, all descriptions apply to **input**  
It's actually much easier to use than **output**  
But it is rather oriented to **form-filling**

More on **flexible** and **free-format** input later

# Integer Descriptors

**I**n (i.e. letter **i**) displays in **decimal**

**Right-justified** in a field of width **n**

**In.m** displays at least **m** digits

**WRITE (\*, '( I7 )')** 123       $\Rightarrow$       ‘      123’

**WRITE (\*, '( I7.5 )')** 123       $\Rightarrow$       ‘      00123’

You can replace the **I** by **B**, **O** and **Z**

For **binary**, **octal** and **hexadecimal**

# Example

```
WRITE (*, '( I7, I7 )') 123, -123
```

```
WRITE (*, '( I7.5, I7.5 )') 123, -123
```

```
    123      -123  
00123    -00123
```

```
WRITE (*, '( B10, B15.10 )') 123, 123
```

```
WRITE (*, '( O7, O7.5 )') 123, 123
```

```
WRITE (*, '( Z7, Z7.5 )') 123, 123
```

```
1111011      0001111011  
 173      00173  
 7B      0007B
```

# Values Too Large

This is **field overflow** on **output**

The **whole field** is replaced by **asterisks**

Putting **1234** into **i4** gives **1234**

Putting **12345** into **i4** gives **\*\*\*\*\***

Putting **-123** into **i4** gives **-123**

Putting **-1234** into **i4** gives **\*\*\*\*\***

This applies to **all numeric** descriptors

Both **REAL** and **INTEGER**

# Fixed-Format REAL

**Fn.m** displays to **m** decimal places

**Right-justified** in a field of width **n**

**WRITE (\*, '( F9.3 )')** 1.23  $\Rightarrow$  ' 1.230'

**WRITE (\*, '( F9.5 )')** 0.123e-4  $\Rightarrow$  ' 0.00001'

You may assume correct **rounding**

Not required, but **traditional** in Fortran

- Compilers may round **exact halves** differently



# Widths of Zero

For **output** a width of **zero** may be used

But only for **formats** **I**, **B**, **O**, **Z** and **F**

It prints the value without any leading spaces

```
write (*, '("/",i0,"/",f0.3)') 12345, 987.654321
```

Prints

```
/12345/987.65
```

# Exponential Format (1)

There are four descriptors: **E**, **ES**, **EN** and **D**  
With the forms **En.m**, **ESn.m**, **ENn.m** and **Dn.m**

All of them use **m** digits after the decimal point  
**Right-justified** in a field of width **n**

**D** is historical – you should avoid it  
Largely equivalent to **E**, but displays **D**

For now, just use **ESn.m** – more on this later

# Exponential Format (2)

The details are complicated and messy  
You don't usually need to know them in detail  
Here are the **two** basic rules for safety

- In **En.m** and **ESn.m**, make  $n \geq m+7$   
That's a good rule for other languages, too
- **Very large** or **small** exponents display oddly  
I.e. exponents outside the range **-99** to **+99**  
Reread using **Fortran formatted input only**

# Numeric Input

F, E, ES, EN and D are similar

- You should use only **Fn.0** (e.g. **F8.0**)  
For extremely complicated reasons
- Any **reasonable** format of value is accepted

There are more details given later

# CHARACTER Descriptor

**An** displays in a **field** with width **n**

Plain **A** uses the width of the **CHARACTER** item

On **output**, if the **field** is too small:

The **leftmost** characters are used

Otherwise:

The text is **right-justified**

On **input**, if the **variable** is too small:

The **rightmost** characters are used

Otherwise:

The text is **left-justified**

# Output Example

```
WRITE (*,'(a3)') 'a'
```

```
WRITE (*,'(a3)') 'abcdefgh'
```

Will display:

```
  a  
abc
```

# Input Example

```
CHARACTER(LEN=3) :: a
```

```
READ (*,'(a8)') a ;  WRITE (*,'(a)') a
```

```
READ (*,'(a1)') a ;  WRITE (*,'(a)') a
```

With input:

abcdefgh

a

Will display:

fgh

a

# LOGICAL Descriptor

**Ln** displays either **T** or **F**

**Right-justified** in a field of width **n**

On **input**, the following is done

- Any **leading spaces** are ignored

- An optional **decimal point** is ignored

- The **next** char. must be **T** (or **t**) or **F** (or **f**)

- Any remaining characters are ignored

E.g. **‘.true.’** and **‘.false.’** are acceptable



# The G Descriptor

The **G** stands for **generalized**

It has the forms **Gn** or **Gn.m**

It behaves according to the **item type**

**INTEGER** behaves like **In**

**CHARACTER** behaves like **An**

**LOGICAL** behaves like **Ln**

**REAL** behaves like **Fn.m** or **En.m**

depending on the size of the value

The rules for **REAL** are fairly sensible

# Other Types of Descriptor

All of the above are **data edit descriptors**  
Each of them matches an item in the **transfer list**  
As mentioned, they must match its **type**

There are some **other types** of descriptor  
These do **not** match a **transfer list** item  
They are executed, and the **next item** is matched

# Text Literal Descriptor

A **string literal** stands for itself, as text

It is displayed just as it is, for output

It is not allowed in a **FORMAT** for input

Using both **quotes** and **apostrophes** helps

The following are all equivalent

```
WRITE (29, '( "Hello" )')
```

```
WRITE (29, "( 'Hello' )")
```

```
WRITE (29, '( ''Hello'' )')
```

```
WRITE (29, "( ""Hello"" )")
```

# Spacing Descriptor

**X** displays a **single blank** (i.e. a **space**)  
It has no **width**, but may be **repeated**

On **input**, it skips over exactly one character

```
READ (*, '(i1, 3x, i1)') m, n  
WRITE (*, '(i1, x, i1, 4x, a)') m, n, '!'

7PQR9
```

Produces '7 9 !'

# Newline Descriptor (1)

/ displays a single **newline** (in effect)

It has no **width**, but may be **repeated**

It can be used as a **separator** (like a comma)

Only if it has no **repeat count**, of course

```
WRITE (*, '(i1/i1, 2/, a)') 7, 9, '!'
```

7

9

!

## Newline Descriptor (2)

On **input**, it skips the rest of the current line

```
READ (*, '(i1/i1, 2/, i1)') l, m, n
```

```
WRITE (*, '(i1, 1x, i1, 1x, i1)') l, m, n
```

```
1 1 1 1
```

```
2 2 2 2
```

```
3 3 3 3
```

```
4 4 4 4
```

Produces “1 2 4”

# Item-Free FORMATs

You can print **multi-line text** on its own

```
WRITE (*, '("Hello" / "Goodbye")')
```

Hello

Goodbye

And skip as many lines as you like

```
READ (*, '(////)')
```

# Generalising That

That is a special case of a **general rule**  
**FORMAT**s are interpreted as far as possible

```
WRITE (*, '(I5, " cubits", F5.2)') 123
```

123 cubits

This reads **42** and skips the following three lines

```
READ (*, '(I3///)') n
```

**42**



# Complex Numbers

For **list-directed** I/O, these are basic types  
E.g. read and displayed like “(1.23,4.56)”

For **formatted** and **unformatted** I/O  
**COMPLEX** numbers are treated as two **REAL**s  
Like an extra **dimension** of extent **two**

```
COMPLEX :: c = (1.23, 4.56)  
WRITE (*, '(2F5.2,3X,2F5.2)') c, 2.0*c
```

```
1.23 4.56      2.46 9.12
```

# Exceptions and IOSTAT (1)

By default, **I/O exceptions** halt the program  
These include an unexpected **end-of-file**

You trap by providing the **IOSTAT** specifier

```
INTEGER :: ioerr
```

```
OPEN (1, FILE='fred', IOSTAT=ioerr)
```

```
WRITE (1, IOSTAT=ioerr) array
```

```
CLOSE (1, IOSTAT=ioerr)
```

# Exceptions and IOSTAT (2)

**IOSTAT** specifies an **integer** variable

**Zero** means **success**, or no detected error

**Positive** means some sort of **I/O error**

An implementation **should** describe the codes

**Negative** means **end-of-file** (but **see later**)

**Fortran 2003** provides its value – see next lecture

# What Is Trapped? (1)

The following are **NOT** errors  
Fortran defines all of this behaviour

- Formatted **READ** beyond **end-of-record**  
Padded with **spaces** to match the **format**

**Fortran 2003** allows a little control of that

- **Writing** a value too large for a **numeric** field  
The **whole field** is filled with **asterisks** (**\*\*\*\*\***)

# What Is Trapped? (2)

The following are **NOT** errors

- **Writing** too long a **CHARACTER** string  
The **leftmost** characters are used
- **Reading** too much **CHARACTER** data  
The **rightmost** characters are used

# What Is Trapped? (3)

The following is what you can **usually** rely on

- End-of-file
- Unformatted **READ** beyond end-of-record  
In most compilers, **IOSTAT** will be negative
- Most **format errors** (syntactically bad values)  
E.g. **12t8** being read as an **integer**

That is roughly the same as **C** and **C++**

# What Is Trapped? (4)

The following are **sometimes** trapped  
The same applies to most other languages

- Numeric overflow (**integer** or **floating-point**)  
**Floating-point** overflow may just deliver **infinity**  
**Integer** overflow may wrap **modulo  $2^N$**   
Or there may be even **less helpful** effects
- ‘Real’ (**hardware** or **system**) I/O errors  
E.g. no space on writing, file server crashing  
**Anything** may happen, and chaos is normal

# 2 GB Warning

I said “**chaos is normal**” and meant it  
Be careful when using files of more than **2 GB**

Most **filesystems** nowadays will support such files  
But not all of the **interfaces** to them do  
Things like **pipes** and **sockets** are different again

- Has **nothing** to do with the Fortran language

Different **compilers** may use different **interfaces**  
And there may be **options** you have to specify



# OPEN

Files are connected to units using OPEN

```
OPEN (UNIT=11, FILE='fred', IOSTAT=ioerr)
```

That will open a sequential, formatted file

You can then use it for either input or output

You can do better, using optional specifiers

Other types of file always need one or more

# Choice of Unit Number

Unit numbers are non-negative integer values

The valid range is system-dependent

You can usually assume that 1–99 are safe

Some may be in use (e.g. for `stdin` and `stdout`)

They are often (not always) 5 and 6

It is simplest to use unit numbers 10–99

Most codes just do that, and have little trouble

- Better ways of doing it covered in next lecture

# ACCESS and FORM Specifiers

These specify the type of I/O and file

‘sequential’ (default) or ‘direct’

‘formatted’ (default) or ‘unformatted’

```
OPEN (UNIT=11, FILE='fred', ACCESS='direct', &  
      FORM='unformatted', RECL=500, IOSTAT=ioerr)
```

That will open a direct-access, unformatted file  
with a record length of 500

You can then use it for either input or output

# Scratch Files

```
OPEN (UNIT=11, STATUS='scratch', &  
      FORM='unformatted', IOSTAT=ioerr)
```

That will open a **scratch** (temporary) file  
It will be **deleted** when it is **closed**

It will be **sequential** and **unformatted**  
That is the most common type of **scratch** file  
But all other types and **specifiers** are allowed

- Except for the **FILE** specifier

# The ACTION Specifier

- This isn't needed, but is **strongly advised**

It helps to protect against mistakes

It enables the reading of **read-only** files

```
OPEN (UNIT=11, FILE='fred', ACTION='read', &  
      IOSTAT=ioerr)
```

Also **'write'**, useful for **pure output** files

The **default**, **'readwrite'**, allows both

# Example (1)

Opening a **text** file for reading data from

```
OPEN (UNIT=11, FILE='fred', ACTION='read', &  
      IOSTAT=ioerr)
```

Opening a **text** file for writing data or results to

```
OPEN (UNIT=22, FILE='fred', ACTION='write', &  
      IOSTAT=ioerr)
```

```
OPEN (UNIT=33, FILE='fred', ACTION='write', &  
      RECL=80, DELIM='quote', IOSTAT=ioerr)
```

## Example (2)

Opening an **unformatted** file for reading from

```
OPEN (UNIT=11, FILE='fred', ACTION='read', &  
      FORM='unformatted', IOSTAT=ioerr)
```

Opening an **unformatted** file for writing to

```
OPEN (UNIT=22, FILE='fred', ACTION='write', &  
      FORM='unformatted', IOSTAT=ioerr)
```

## Example (3)

Opening an **unformatted** workspace file  
It is your choice whether it is temporary

```
OPEN (UNIT=22, STATUS='scratch', &  
      FORM='unformatted', IOSTAT=ioerr)
```

```
OPEN (UNIT=11, FILE='/tmp/fred', &  
      FORM='unformatted', IOSTAT=ioerr)
```

See extra slides for **direct-access** examples



# Omitted For Sanity

These are in the extra slides

Techniques for reading **free-format** data

Some more detail on **formatted** I/O

**Internal files** and **dynamic formats**

More on **OPEN**, **CLOSE**, **positioning** etc.

**Direct-access** I/O

There are extra, extra slides on some details

# Introduction to Modern Fortran

## *Data Pointers*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Data Pointers

- Fortran pointers are unlike C/C++ ones  
Not like Lisp or Python ones, either
  - Errors with using pointers are rarely obvious  
This statement applies to almost all languages
  - Fortran uses a semi-safe pointer model  
Translation: your footgun has a trigger guard
- Use pointers only when you need to

# Pointer and Allocatable

**Pointers** are a sort of changeable **allocation**  
In that use, they almost always point to **arrays**  
For example, needed for **non-rectangular arrays**

**Always** try to use **allocatable** arrays first  
Only if they really aren't adequate, use **pointers**

**ALLOCATABLE** was restricted in **Fortran 95**  
**Fortran 2003** removed almost all **restrictions**  
You may come across **POINTER** in old code  
It can usually be replaced by **ALLOCATABLE**

# Pointer-Based Algorithms

Some genuinely **pointer-based algorithms**

Fortran is not really ideal for such uses

- But don't assume anything else is any better!

There are **NO** safe pointer-based languages

Theoretically, one could be designed, but ...

In Fortran, see if you can use **integer indices**

That has **software engineering** advantages, too

If you can't, you may have to use **pointers**

# Pointer Concepts

**Pointer** variables point to **target** variables

In almost all uses, **pointers** are **transparent**

- You access the **target variables** they point to

**Dereferencing** the **pointer** is **automatic**

- Special syntax for meaning the **pointer value**

The **POINTER** **attribute** indicates a **pointer**

The **TARGET** **attribute** indicates a **target**

No variable can have both **attributes**

# Example

```
PROGRAM fred
  REAL, TARGET :: popinjay = 0.0
  REAL, POINTER :: arrow
  arrow => popinjay
  ! arrow now points to popinjay
  arrow = 1.23
  PRINT *, popinjay
  popinjay = 4.56
  PRINT *, arrow
END PROGRAM fred
```

1.2300000

4.5599999

# Pointers and Target Arrays

```
REAL, DIMENSION(20), TARGET :: array  
REAL, DIMENSION(:), POINTER :: index
```

Pointer arrays must be declared without bounds  
They will take their bounds from their targets

- Pointer arrays have just a rank  
Which must match their targets, of course

Very like allocatable arrays



# Use of Targets

Treat **targets** just like ordinary **variables**

The **ONLY** difference is an extra **attribute**  
Allows them on the **RHS** of **pointer assignment**

Valid **targets** in a **pointer assignment**?

If OK for **INTENT(INOUT)** actual argument  
**Variables**, **array elements**, **array sections** etc.

```
REAL, DIMENSION(20, 20), TARGET :: array  
REAL, DIMENSION(:, :), POINTER :: index  
index => array(3:7:2, 8:2:-1)
```

# Initialising Pointers

Pointer variables are initially undefined

- Not initialising them is a **Bad Idea**
- You can use the special syntax `=> null()`  
To initialise them to **disassociated** (*sic*)

```
REAL, POINTER :: index => null()
```

- Or you can point them at a **target**, **ASAP**  
Note that `null()` is a **disassociated** target

# Pointer Assignment

You use the special **assignment operator** **=>**

Note that using **=** assigns to the **target**

```
PROGRAM fred
  REAL, TARGET :: popinjay
  REAL, POINTER :: arrow
  arrow => popinjay      ! POINTER assignment
  ! arrow now points to popinjay
  arrow = 1.23           ! TARGET assignment
  PRINT *, popinjay
  popinjay = 4.56        ! TARGET assignment
  PRINT *, arrow
  arrow => null()         ! POINTER assignment
END PROGRAM fred
```

# Pointer Expressions

Also **pointer expressions** on the RHS of **=>**  
Currently, only the **results of function calls**

```
FUNCTION select (switch, left, right)
  REAL, POINTER :: select, left, right
  LOGICAL switch
  IF (switch) THEN
    select => left
  ELSE
    select => right
  END IF
END FUNCTION select
```

```
new_arrow => select(A > B, old_arrow, null())
```

# ALLOCATE

You can use this just as for **allocatable arrays**  
This creates some space and sets up **array**

```
REAL, DIMENSION(:, :), POINTER :: array  
ALLOCATE(array(3:7:2, 8:2:-1), STAT=n)
```

If you can, stick to using **ALLOCATABLE**

Do you get the idea I don't like pointers much?  
At the end, I mention why you may need them

# DEALLOCATE

- Only on **pointers** set up by **ALLOCATE**

**DEALLOCATE**(array, STAT=n)

**array** now becomes **disassociated**

**Other** pointers to its target become **undefined**

- Don't **DEALLOCATE** **undefined** pointers

That is **undefined** behaviour

# Previous Pointer Values

New **pointer value** overwrites the previous one  
Applies to both **assignment** and **ALLOCATE**  
Well, it is a sort of **assignment** ...

- Does not affect **other pointers** to the **target**

But **DEALLOCATE** makes other pointers **undefined**  
Also happens if the **target** goes out of **scope**

- That causes the **dangling pointer** problem

And **assignment** can break the last **link**

- **Memory leaks** and (**rarely**) worse problems

# ASSOCIATED

- Can test if **pointers** are **associated**

```
IF (ASSOCIATED(array)) . . .  
IF (ASSOCIATED(array, target)) . . .
```

Works if **array** is **associated** or **disassociated**

Latter tests if **array** is **associated** with **target**

- Don't use it on **undefined pointers**

That is **undefined** behaviour



# A Nasty “Gotcha”

Fortran 95 forbids **POINTER** and **INTENT**

- Fortran 2003 applies **INTENT** to the **link**

```
subroutine joe (arg)
  real, target :: junk
  real, pointer, intent(in) :: arg
  allocate(arg)      ! this is ILLEGAL
  arg => junk         ! this is ILLEGAL
  arg = 4.56          ! but this is LEGAL :-(
end subroutine joe
```

# Irregular Arrays

- Fortran does not support them

This is how you do the task, if you need to

```
TYPE Cell
```

```
    REAL, DIMENSION(:), ALLOCATABLE :: column  
END TYPE Cell
```

```
TYPE(Cell), DIMENSION(:), ALLOCATABLE :: matrix
```

`matrix` can be a non-rectangular matrix

Note that `pointers` are not needed in this case

# Example

```
TYPE Cell
    REAL, DIMENSION(:), ALLOCATABLE :: column
END TYPE Cell

TYPE(Cell), DIMENSION(:), ALLOCATABLE :: matrix

INTEGER, DIMENSION(100) :: rows
READ *, N, (rows(K), K = 1,N)
ALLOCATE(matrix(1:N))
DO K = 1,N
    ALLOCATE(matrix(K)%column(1:rows(K)))
END DO
```

# Arrays of Pointers

- Fortran does not support them

This is how you do the task, if you need to

```
TYPE Cell
```

```
    REAL, DIMENSION(:), POINTER :: column  
END TYPE Cell
```

```
TYPE(Cell), DIMENSION(100) :: matrix
```

# Remember Trees?

This was the example we used in **derived types**

```
TYPE :: Leaf
    CHARACTER(LEN=20) :: name
    REAL(KIND=dp), DIMENSION(3) :: data
END TYPE Leaf
TYPE :: Branch
    TYPE(Leaf), ALLOCATABLE :: leaves(:)
END TYPE Branch
TYPE :: Trunk
    TYPE(Branch), ALLOCATABLE :: branches(:)
END TYPE Trunk
```

# Recursive Types

We can do this more easily using **recursive types**

```
TYPE :: Node
    TYPE(Node), POINTER :: subnodes(:)
    CHARACTER(LEN=20) :: name
    REAL(KIND=dp), DIMENSION(3) :: data
END TYPE Node
```

Recursive components must be **pointers**

**Fortran 2008** will allow **allocatable**

Obviously a type cannot include itself directly

# More Complicated Structures

In mathematics, a **graph** is a set of **linked nodes**  
Common forms include **linked lists**, **trees** etc.





A **tree** is just a hierarchy of objects  
We have already covered these, in principle

**Linked lists** (also called **chains**) are common  
And there are lots of more complicated structures

Those are very painful to handle in old Fortran  
So most Fortran programmers tend to avoid them  
But they aren't difficult in modern Fortran

# Singly Linked List

Head







Tail



# Doubly Linked List

Head



Tail

# Linked Lists

You can handle **linked lists** in a similar way  
And any other graph-theoretic data structure, too

```
TYPE Cell
  CHARACTER(LEN=20) :: node_name
  REAL :: node_weight
  TYPE(Cell), POINTER :: next, last, &
    first_child, last_child
END TYPE Cell
```

Working with such data structures is non-trivial  
Whether in Fortran or any other language

# Graph Structures

Using pointers in **Fortran** is somewhat tedious  
But it is as easy as in **C++** and a **little** safer

**Graph structures** are in **computer science**  
**linked lists** are probably the only easy case  
Plenty of books on them, for example:

**Cormen, T.H. et al. Introduction to Algorithms**  
**Knuth, D.E. The Art Of Computer Programming**  
Also **Sedgewick, Ralston, Aho et al.** etc.

# Procedure Pointers

Fortran 2003 allows them, as well as data pointers

Don't go there

This has absolutely **nothing** to do with Fortran

They are a nightmare in all languages, including C++

They are **almost impossible** to use safely

A **fundamental** problem in any **scoped** language

- **Very rarely** need them in **clean code**, anyway

Passing **procedures as arguments** is usually enough

Or one **procedure** calling a **fixed set** of others

# Introduction to Modern Fortran

## *Advanced Array Concepts*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Summary

This will describe some advanced array features  
Use them only when you need their facilities

It will also cover some aspects of array use  
Important for correctness and performance

There is a lot more on both

- Please ask if you need any help

# Testing Allocation

Can test if an **ALLOCATABLE** object is **allocated**

```
INTEGER, DIMENSION(:), ALLOCATABLE :: counts
```

```
• • •  
IF (ALLOCATED(counts)) THEN  
  • • •
```

**Warning:** rules of (de)allocation are non-trivial

Can happen **automatically** under some circumstances

Generally, restructure your code to **not need it**

# Higher Rank Constructors

Constructors create only **rank one** arrays

We shall now see how to construct higher ranks

It is done by constructing a **rank one** array

And then mapped using the **RESHAPE** function

This is very easy, but looks a bit messy



# The RESHAPE Intrinsic (1)

This allows arbitrary restructuring of arrays  
The following is only its **very** simplest use

RESHAPE (source, shape)

**source** provides the data in array element order  
**shape** specifies the **shape** of array to deliver

# The RESHAPE Intrinsic (2)

```
REAL, DIMENSION(3, 4) :: array
```

```
array = RESHAPE( (/ 1.1, 2.1, 3.1, 1.2, 2.2, &  
                  3.2, 1.3, 2.3, 3.3, 1.4, 2.4, 3.4 /), (/ 3, 4 /) )
```

Is functionally equivalent to:

```
DO m = 1, 3  
  DO n = 1, 4  
    array(m, n) = m+0.1*n  
  END DO  
END DO
```

# The RESHAPE Intrinsic (3)

It can be used in **constant expressions**

```
REAL, DIMENSION(3, 4) :: array = &  
    RESHAPE( (/ 1.1, 2.1, 3.1, 1.2, 2.2, &  
              3.2, 1.3, 2.3, 3.3, 1.4, 2.4, 3.4 /), (/ 3, 4 /) )
```

It also allows **arbitrary reordering**

And **padding** with copies of an array

See the references for more details

# Example

Create the zero vector, and the three unit vectors

```
REAL, DIMENSION(1:3), PARAMETER :: &  
    vec_0 = (/ 0.0, 0.0, 0.0 /), &  
    vec_i = (/ 1.0, 0.0, 0.0 /), &  
    vec_j = (/ 0.0, 1.0, 0.0 /), &  
    vec_k = (/ 0.0, 0.0, 1.0 /)
```

Create the identity matrix

```
REAL, DIMENSION(1:3, 1:3), PARAMETER :: &  
    identity = RESHAPE( (/ vec_i, vec_j, vec_k /), (/ 3, 3 /) )
```

# RESHAPE More Generally

It isn't restricted to **multi-dim. constants**

You can use it for fancy array restructuring

- Study the specification before doing that  
Restructuring arrays is dangerous territory
- And there are several other such intrinsics  
I.e. ones with **important** uses but no **simple** uses

# Vector Indexing (1)

Vectors may be used as indices

```
INTEGER, DIMENSION(1:5) :: &  
    j = (/ 3, 1, 5, 2, 4 /), k = (/ 2, 3, 2, 1, 3 /)  
REAL, DIMENSION(1:5) :: x, &  
    y = (/ 1.2, 2.3, 3.4, 4.5, 5.6 /)  
x(j) = y(k)  
PRINT *, y(k)  
PRINT *, x
```

```
2.3000000  3.4000001  2.3000000  1.2000000  3.4000001  
3.4000001  1.2000000  2.3000000  3.4000001  2.3000000
```

# Vector Indexing (2)

Using **vector indices** is a bit like **sections**  
There are important differences – be careful

You can them for **reading** arrays quite safely  
Elements must be **distinct** if **updating**

- **NOT** recommended for use in **arguments**  
If used in arguments, those **must not** be updated  
And it forces the compiler to **copy** the array

# Masked Assignment (1)

Set all negative values in an array A to zero

```
REAL, DIMENSION(20, 30) :: array
```

```
DO j = 1,30  
    DO k = 1,20  
        IF (array(k,j) < 0.0) array(k,j) = 0.0  
    END DO  
END DO
```

But the **WHERE** statement is more convenient

```
WHERE (array < 0.0) array = 0.0
```



# Masked Assignment (2)

It has a **statement construct** form, too

```
WHERE (array < 0.0)
    array = 0.0
ELSE WHERE
    array = 0.01*array
END WHERE
```

**Masking expressions** are **LOGICAL** arrays

You can use an actual array there, if you want

**Masks** and **assignments** need the same **shape**

# Masked Assignment (3)

Fortran 2003 extends it considerably

Don't use LHS arrays in **non-elemental** functions

The following is asking for trouble:

```
WHERE (arr1 < arr2)
    arr1 = 1.0
ELSE WHERE
    arr2 = sum(arr1)
END WHERE
```

- Don't bother with the **FORALL** statement

# Memory Efficiency (1)

**Local** arrays can be implemented in many ways  
Only a few **Ada** compilers handle them properly

You can exhaust your program's **stack** with them  
**Too big**, or **too many** due to **deep recursion**

- It will usually cause a **truly** horrible crash

**Allocatable** arrays always go on the '**heap**'

**Automatic** arrays **often** go on the '**heap**'

That is less efficient, but is handled much better

- Making all big arrays **allocatable** isn't stupid

# Memory Efficiency (2)

As always, every solution has its own problems  
Lots of **allocation** and **deallocation** isn't ideal

- Each **(de)allocation** costs some CPU time  
Not generally a problem for Fortran programs
- Poor compilers may have **memory leaks**  
Most Fortran compilers don't have them badly

Both **because of** the language's **restrictions**

# Memory Efficiency (3)

- The big problem is memory fragmentation  
Describing how and why is beyond this course  
Luckily, in AD 2007, there is a simple solution
- Best one is to use 64-bit addressing  
Gets rid of the worst of the problems, painlessly  
I do that, even on systems with 2 GB of memory
- Please ask if you want to know more

# Order of Evaluation (1)

Array assignments etc. are like implicit loops

But, except in I/O, no order of evaluation implied

Also the behaviour is different when modifying

- Each pass of a loop is executed in order
- Array assignments do it all “in parallel”
- You should avoid code where it matters

The compiler may have to copy the array

It risks confusion when tuning your code

# Order of Evaluation (2)

```
INTEGER, DIMENSION(5) :: array = (/ 1, 2, 3, 4, 5 /)
array(2:5) = array(1:4)
PRINT *, array
```

```
array = (/ 1, 2, 3, 4, 5 /)
DO k = 1,4
    array(k+1) = array(k)
END DO
PRINT *, array
```

```
1  1  2  3  4
1  1  1  1  1
```

# Performance (1)

- Efficient use of arrays is critical  
This course has **NOT** taught any of that  
It covers quite enough without adding it!
- **Generally**, follow this procedure:

Start by writing **clean and clear** code  
Get it working, and test it fairly thoroughly  
If too slow, use a **profiler** to see where  
And **only then** tune **only those** aspects



# Performance (2)

You get **most gain** by using **faster methods**  
**Followed** by the following aspects:

- Improve the **layout** and **access patterns**  
This is **locality** (improved cache usage etc.)

- Avoid unnecessary **array copying**  
Compilers often have to do that for some codes  
Some compilers copy when they **don't** need to

- Improve the actual CPU efficiency  
This is getting into advanced tuning

# Memory Locality (1)

Things **used together** should be **stored together**  
Remember that “**first index varies fastest**”

```
REAL, DIMENSION(3000, 5000) :: array
DO n = 1, 5000
    DO m = 1, 3000
        array(m, n) = m+0.1*n
    END DO
END DO
```

- Note that the first index varies fastest

# Memory Locality (2)

Sections and masking can cause trouble

```
REAL, DIMENSION(1000, 1000) :: array  
CALL FRED( array(123, :) )
```

The elements of the vector are a long way apart  
A problem if **FRED** accesses it a lot

- Consider making a temporary copy of it

# Access Patterns

- **Sequential access** is generally efficient  
Avoid non-sequential access wherever possible
- This can be **much** slower than sequential

```
REAL, DIMENSION(1000) :: arr1, arr2  
INTEGER, DIMENSION(1000) :: random  
arr1(random) = arr2(random)
```

# Unnecessary Copying (1)

It is hard to describe when this may occur  
It helps if you can mentally compile the code

- Avoiding using the **LHS** array on the **RHS**  
Except when the uses are purely **elemental**
- Generally, **sections** do not need a copy  
Unlike **arguments** with **vector indexed arrays**
- Compilers often do **unnecessary** copying  
In a very bad case, even for **CALL Fred(data(:))**

# Example

```
INTEGER :: arr1(1:50), arr2(1:100), arr3(1:100)  
REAL, DIMENSION(20, 20) :: mat1, mat2, mat3
```

These shouldn't require a copy

```
arr1 = arr1+arr2(1:50)+arr3(arr2(51:100))  
mat1 = MATMUL(mat2, mat3)
```

But these almost certainly will

```
arr1 = arr1(:, -1)+arr2(1:50)  
mat1 = MATMUL(mat1, mat2)
```

# Unnecessary Copying (2)

And, while this **shouldn't**, ...

```
mat1 = mat1 + MATMUL(mat2, mat3)
```

There is more on this under **procedures**

- Generally, don't worry unless you have to  
If your program runs fast enough, who cares?
- If not, **time** and **profile** it first  
Ask for advice if you have problems

# High-Performance Problems

There are some other problems some people hit  
Too complicated to even describe here

- Ignore them **until** you have problems  
Then ask for help with tackling them

Buzzwords and phrases include:

TLB thrashing

Cache conflicts

False sharing

Memory banking



# Reminder

- You don't have to remember all of this
  - Start by using the simplest features only
  - Use the fancy ones only when you need them
- If you know they exist, you can look them up

# Introduction to Modern Fortran

## *Advanced Use Of Procedures*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Summary

We have omitted some important **concepts**  
They are complicated and confusing

There are a lot of **features** we have omitted  
Mostly because they are **hard to use** correctly  
And sometimes because they are **inefficient**

This lecture covers some of the most **important**

- Refer to this when you need to

# ALLOCATABLE and POINTER

You can pass **ALLOCATABLE** and **POINTER** arrays

In the usual case, the **procedure** has **neither**

The dummy argument is **associated** with the data

- You can't reallocate or redirect **in the procedure**

To do that, declare the **dummy argument** as

**ALLOCATABLE** or **POINTER**, as appropriate

**Warning** for **INTENT(OUT)** and **ALLOCATABLE**:

These are **deallocated** on entry, even if not used

# Association (1)

Fortran uses **argument association** in calls  
**Dummy** arguments **refer to** the **actual** ones

- You don't need to know exactly how it is done  
It may be **aliasing** or **copy-in/copy-out**

**Expressions** are stored in a hidden **variable**

The **dummy** argument is **associated** with that

- It obviously must not be updated in any way

Using **INTENT** is **strongly** recommended

# Association (2)

```
REAL, DIMENSION(1:10, 1:20, 1:3) :: data  
CALL Fred (data(:, 5:15, 2), 1.23*xyz )
```

```
SUBROUTINE Fred (array, value)  
REAL, DIMENSION(:, :) :: array  
REAL, INTENT(IN) :: value
```

array in fred refers to data(:, 5:15, 2)

value refers to a location containing 1.23\*xyz

# Updating Arguments (1)

A **dummy argument** must not be updated if:

- The **actual argument** is an **expression**
- It **overlaps** another argument in any way

```
REAL, DIMENSION(1:20, 1:3) :: data  
CALL Fred (data(5:15, 2), data(17:, 2))
```

```
SUBROUTINE Fred (arr1, arr2)  
REAL, DIMENSION(:) :: arr1, arr2  
arr1 = 1.23 ;   arr2 = 4.56
```

- The above works as you expect

# Updating Arguments (2)

```
REAL, DIMENSION(1:20, 1:3) :: data  
CALL Fred (data(5:15, 2), data(1:10, 2))
```

```
SUBROUTINE Fred (arr1, arr2)  
REAL, DIMENSION(:) :: arr1, arr2  
arr2(1, 1) = 4.56
```

- The above is **not** allowed  
Because **arr1** and **arr2** overlap

Even though **arr2(1, 1)** is not part of **arr1**



# Updating Arguments (3)

```
REAL :: X  
CALL Fred (X + 0.0)
```

```
SUBROUTINE Fred (Y)  
Y = 4.56
```

- The above is **not** allowed – obviously
- That also applies to **array expressions**  
**Vector indexing** behaves like an **expression**

# Warning for C/C++ People

```
REAL, DIMENSION(1:20) :: data  
CALL Fred (data(2), data)
```

```
SUBROUTINE Fred (var, array)  
REAL :: var  
REAL, DIMENSION(:) :: array  
array = 4.56
```

- The above is **not** allowed, either

Even **array elements** are **associated**

# Using Functions

**Functions** are called just like built-in ones  
They may be **optimised** in similar ways

```
REAL :: scale, data(1000)
```

```
...
```

```
READ *, scale ! assume that this reads 0.0
```

```
Z = Variance(data)/(scale+Variance(data))
```

**Variance** may be called **0**, **1** or **2** times

# Impure Functions

Pure functions have defined behaviour

- Whether they are declared **PURE** or not

Impure functions occasionally misbehave

Generally, because they are over-optimised

There are rules for **safety in practice**

But they are too complicated for this course

- Ask if you need help with this

# FUNCTION Result Variable

The **function name** defines the **result variable**  
You can change this if you prefer

```
FUNCTION Variance_of_an_array (Array) RESULT(var)
  REAL :: var
  REAL, INTENT(IN), DIMENSION(:) :: Array
  var = SUM(Array)/SIZE(Array)
  var = SUM((Array-var)**2)/SIZE(Array)
END FUNCTION Variance_of_an_array

REAL, DIMENSION(1000) :: data
...
Z = Variance_of_an_array(data)
```

# PURE Subroutines

You can declare a **subroutine** to be **PURE**

Like **functions**, but with one **fewer** restriction  
**INTENT(OUT)** and **INTENT(INOUT)** are allowed

```
PURE SUBROUTINE Init (array, value)
    REAL, DIMENSION(:), INTENT(OUT) :: array
    REAL, INTENT(IN) :: value
    array = value
END SUBROUTINE Init
```

They can be declared as **ELEMENTAL**, too

# Recursion

Fortran 90 allowed this for the first time

**Recursive** procedures must be declared as such

- If you don't, **recursion** **may** cause chaos

```
RECURSIVE SUBROUTINE Chop (array, value)
```

```
...
```

- Avoid it unless you actually **need** it
- Check **all** procedures in the recursive loop

# OPTIONAL Arguments

- Use **OPTIONAL** for setting **defaults** only  
On entry, check and copy **ALL** args  
Use **ONLY** local copies thereafter  
Now, all variables are well defined when used
- Can do the converse for optional **results**  
Just before returning, check and copy back
- Beyond this should be done only by **experts**



# OPTIONAL Example (1)

```
FUNCTION fred (alf, bert)
  REAL :: fred, alf, mybert
  REAL, OPTIONAL, INTENT(IN) :: bert
  IF (PRESENT(bert)) THEN
    mybert = bert
  ELSE
    mybert = 0.0
  END IF
```

Now use **mybert** in rest of procedure

## OPTIONAL Example (2)

```
SUBROUTINE fred (alf, bert)
REAL :: alf
REAL, OPTIONAL, INTENT(OUT) :: bert
...
IF (PRESENT(bert)) bert = ...
END SUBROUTINE fred
```

# Fortran 2003

Adds potentially useful **VALUE** attribute

See **OldFortran** course for information

Seriously. It's also useful for **conversion**

And the **PROCEDURE** declaration statement

A cleaner and more modern form of **EXTERNAL**

Its **usage** is not what you would expect, though

And probably more ...

# Arrays and CHARACTER

We have **over-simplified** these so far

No problem, if you use only recommended style

- You need to know more if you go **beyond** that
- We start by describing what you **can** do  
Including some warnings about **efficient** use

And then continue with how it actually works

# Array Valued Functions

Arrays are **first-class objects** in Fortran  
Functions can return array results

- In practice, doing so always needs a **copy**  
However, don't worry too much about this

Declare the **function** just as for an **argument**  
The constraints on the **shape** are similar

- If it is too slow, ask for advice

# Example

This is a bit futile, but shows what can be done

```
FUNCTION operate (mat1, mat2, mat3)
  IMPLICIT NONE
  REAL, DIMENSION(:, :), INTENT(IN) :: &
    mat1, mat2, mat3
  REAL, DIMENSION(UBOUND(mat1, 1), &
    UBOUND(mat2, 2)) :: operate
  ! Checking omitted, again
  operate = MATMUL(mat1, mat2) + mat3
END FUNCTION operate
```

# Array Functions and Copying

The **result** need not be **copied** on return

The **interface** provides enough information

In practice, don't bet on it ...

**Array functions** can also **fragment memory**

Ask if you want to know how and why

- Generally a problem **only** for **HPC**  
I.e. when either **time** or **memory** are bottlenecks

# What Can Be Done

- Just use **array functions** regardless  
If you don't have a problem, why worry?
- **Time** and **profile** your program  
Tune only code that is a **bottleneck**
- Rewrite **array functions** as **subroutines**  
I.e. turn the **result** into an **argument**
- Use **ALLOCATABLE** results (sic)
- Ask for further advice with tuning



# CHARACTER And Copying

In this respect, CHARACTER  $\equiv$  array

Most remarks about arrays apply, unchanged

- But it is only rarely important

Fortran is rarely used for heavy character work

It works fairly well, but it isn't ideally suited

Most people find it very tedious for that

- If you need to, ask for advice

# Character Valued Functions (1)

Earlier, we considered just one form  
Almost anything more needs a **copy**  
Some compilers will copy even those

- Often, the cost of that does not matter

You are not restricted to just that form  
Declare the **function** just as for an **argument**  
The constraints on the **shape** are similar

- If it is too slow, ask for advice

# Character Valued Functions (2)

The **result length** can be taken from an **argument**

```
FUNCTION reverse_word (word)
  IMPLICIT NONE
  CHARACTER(LEN=*), INTENT(IN) :: word
  CHARACTER(LEN=LEN(word)) :: reverse_word
  INTEGER :: I, N
  N = LEN(word)
  DO I = 1, N
    reverse_word(I:I) = word(N+1-I:N+1-I)
  END DO
END FUNCTION reverse_word
```

# Character Valued Functions (3)

This is a bit futile, but shows what can be done  
The **result length** is a non-trivial **expression**

```
FUNCTION Interleave (text1, count, text2)
  IMPLICIT NONE
  CHARACTER(LEN=*), INTENT(IN) :: text1, text2
  INTEGER, INTENT(IN) :: count
  CHARACTER(LEN=LEN(text1)+count+ &
    LEN(text2)) :: Interleave
  Interleave = text1 // REPEAT(' ', count) // text2
END FUNCTION Interleave
```

# Explicit/Assumed Size/Shape (1)

- The good news is that everything works  
Can mix **assumed** and **explicit** *ad lib*.

There are some potential performance problems

- Passing **assumed** to **explicit** forces a **copy**
- It can be a problem calling some libraries  
Especially ones written in **old Fortran**
- Write clean code, and see if it is fast enough  
If you find that it isn't, ask for advice

# Explicit/Assumed Size/Shape (2)

This code is not a problem:

```
SUBROUTINE Weeble (matrix)
    REAL, DIMENSION(:, :) :: matrix
END SUBROUTINE Weeble
```

```
SUBROUTINE Burble (space, M, N)
    REAL, DIMENSION(M, N) :: space
    CALL Weeble(space)
END SUBROUTINE Burble
```

```
REAL, DIMENSION(100,200) :: work
CALL Burble(work, 100, 200)
```

# Explicit/Assumed Size/Shape (3)

Nor even something as extreme as this:

```
SUBROUTINE Weeble (matrix)
    REAL, DIMENSION(:, :) :: matrix
END SUBROUTINE Weeble
```

```
SUBROUTINE Burble (space, N, J1, K1, J2, K2)
    REAL, DIMENSION(N, *) :: space
    CALL Weeble(space(J1:K1, J2:K2))
END SUBROUTINE Burble
```

```
REAL, DIMENSION(100, 200) :: work
CALL Burble(work, 100, 20, 80, 30, 70)
```

# Explicit/Assumed Size/Shape (4)

But this code forces a copy:

```
SUBROUTINE Bubble (matrix, M, N)
    REAL, DIMENSION(M, N) :: matrix
END SUBROUTINE Bubble
```

```
SUBROUTINE Womble (space)
    REAL, DIMENSION(:, :) :: space
    CALL Bubble(space, UBOUND(space, 1), &
        UBOUND(space, 2))
END SUBROUTINE Womble
```

```
REAL, DIMENSION(100,200) :: work
CALL Womble(work)
```



# Example – Calling LAPACK

LAPACK is written in Fortran 77

It cannot handle assumed shape arrays

So here is how to call SPOTRF (Cholesky)

```
SUBROUTINE Chol (matrix, info)
  REAL, DIMENSION(:, :), INTENT(INOUT) :: matrix
  INTEGER, INTENT(INOUT) :: info
  CALL SPOTRF('L', UBOUND(matrix, 1), &
             matrix, UBOUND(matrix, 1), info)
END SUBROUTINE Chol
```

matrix will be copied on call and return

# Sequence Association (1)

Have covered **assumed shape** and **char. length**  
And **explicit shape** and **char. length**  
but only when the **dummy** and **actual** match

- That constraint is **not required** (nor **checked**)

You need to know an extra concept to go further  
That is called **sequence association**

- You are recommended to go cautiously here  
Don't do it until you are confident with Fortran

# Sequence Association (2)

Explicit shape and assumed size arrays only

If the dummy and actual bounds do not match

Argument is flattened in array element order

And is given a shape by the dummy bounds

Exactly the way the RESHAPE intrinsic works

There are important uses of this technique

- Or you can shoot yourself in the foot

# Example

```
SUBROUTINE operate_1 (vector, N)
    REAL, DIMENSION(N) :: vector
    . . .
SUBROUTINE operate_2 (matrix, M, N)
    REAL, DIMENSION(M, N) :: matrix
    . . .

REAL, DIMENSION(1000000) :: workspace
. . .
IF (cols = 0) THEN
    CALL operate_1(workspace, rows)
ELSE
    CALL operate_2(workspace, rows, cols)
END IF
```

# Sequence Association (3)

The same holds for **explicit length CHARACTER**  
Everything is concatenated and then reshaped

Character lengths are like an extra **dimension**  
Naturally, it varies **faster** than the first **index**

One restriction needed to make this work  
**Assumed shape** arrays of **CHARACTER**  
need **assumed length** or **matching** lengths

# Example

```
SUBROUTINE operate (fields, N)  
    CHARACTER(LEN=8), DIMENSION(10, N) :: fields  
END SUBROUTINE operate
```

```
CHARACTER(LEN=80), DIMENSION(1000) :: lines
```

```
...
```

```
! Read in N lines
```

```
CALL operate(lines, N)
```

# Implicit Interfaces (1)

Calling an undeclared procedure is allowed  
The **actual arguments** define the **interface**

- I strongly recommend not doing this  
Mistyped array names often show up as **link errors**

```
REAL, DIMENSION(1000) :: lines
```

```
...
```

```
lines(5) = lones(7)
```

Undefined symbol lones\_ in file test.o

# Implicit Interfaces (2)

Only Fortran 77 interface features can be used  
The args and result must be exactly right  
Must declare the result type of functions

```
REAL, DIMENSION(KIND=dp) :: DDOT  
...  
X = DDOT(array)
```

- This is commonly done for external libraries  
I.e. ones that are written in Fortran 77, C etc.
- Interface modules are a better way



# EXTERNAL

This declares an **external procedure** name

It's **essential** only when passing as argument

I.e. if the procedure name is **used** but not **called**

- I recommend it for all **undeclared** procedures  
More as a form of documentation than anything else
- But **explicit interfaces** are always better

# Example

Here is the **LAPACK** example again

```
SUBROUTINE Chol (matrix, info)
  REAL, DIMENSION(:, :), INTENT(INOUT) :: matrix
  INTEGER, INTENT(INOUT) :: info
  EXTERNAL :: SPOTRF
  CALL SPOTRF('L', UBOUND(matrix, 1), &
             matrix, UBOUND(matrix, 1), info)
END SUBROUTINE Chol
```

# Introduction to Modern Fortran

## *Advanced I/O and Files*

Nick Maclaren

[nmm1@cam.ac.uk](mailto:nmm1@cam.ac.uk)

March 2014

# Summary

This will describe some advanced I/O features  
Some are useful but only in Fortran 2003  
Some are esoteric or tricky to use

- The points here are quite important  
Excluded only on the grounds of time

There is a lot more in this area

- Please ask if you need any help

# Partial Records in Sequential I/O

Reading only part of a **record** is supported  
Any unread data in the **record** are skipped  
The next **READ** uses the next **record**

Fortran 90 allows you to change that

- But **ONLY** for **formatted**, **external** I/O

Specify **ADVANCE='no'** in the **READ** or **WRITE**  
This is called **non-advancing** I/O

# Non-Advancing Output

You can build up a **record** in sections

```
WRITE (*, '(a)', ADVANCE='no') 'value = '  
IF (value < 0.0) THEN  
    WRITE (*, '("None")') value  
ELSE  
    WRITE (*, '(F5.2)') value  
END IF
```

This is, regrettably, the only portable use

# Use for Prompting

```
WRITE (*, '(a)', ADVANCE='no') 'Type a number: '  
READ (*, *) value
```

That will usually work, but may not

The text may not be written out immediately  
Even using **FLUSH** may not force that

Too many prompts may exceed the **record length**

# Non-Advancing Input

You can decode a **record** in sections  
Just like for **output**, if you know the **format**

Reading **unknown length** records is possible  
Here are two **recipes** that are safe and reliable

Unfortunately, **Fortran 90** and **Fortran 2003** differ



# Recipe (1) - Fortran 2003

```
USE, INTRINSIC :: ISO_FORTRAN_ENV  
CHARACTER, DIMENSION(4096) :: buffer  
INTEGER :: status, count  
READ (1, '(4096a)', ADVANCE='no', SIZE=count, &  
      IOSTAT=status) buffer
```

If **IOSTAT** is **IOSTAT\_EOR**, the record is short

If **IOSTAT** is **IOSTAT\_END**, we are at **end-of-file**

**SIZE** returns the number of **characters** read

## Recipe (2) - Fortran 90

```
CHARACTER, DIMENSION(4096) :: buffer  
INTEGER :: count  
READ (1, '(4096a)', ADVANCE='no', SIZE=count, &  
      EOR=10, EOF=20) buffer
```

The **EOR** branch is taken if the record is short  
The following happens whether or not it is

**SIZE** returns the number of **characters** read

# General Free-Format Input

- Can read in whole lines, as described above  
And then decode using **CHARACTER** operations  
You can also use **internal files** for conversion
- Can use some other language for conversion  
I use **Python**, but **Perl** is fine, too  
Use it to convert to a Fortran-friendly format
- You can call **C** to do the conversion  
That isn't always as easy as people think it is

# List-Directed I/O (1)

This course has massively over-simplified  
All you need to know for simple test programs  
It is used mainly for diagnostics etc.

Here are a few of its extra features

Separation is by **comma**, **spaces** or both  
That is why **comma** needs to be **quoted**  
Theoretically, that can happen on output, too

## List-Directed I/O (2)

You may use **repeat counts** on **values**

**100\*1.23** is a **hundred** repetitions of **1.23**

That is why **asterisk** needs to be **quoted**

Theoretically, that can happen on output, too

There may be **null values** in input

**“1.23 , , 4.56”** is **1.23** , **null value**, **1.234.56**

**“100\* ”** is a **hundred null values**

Null values suppress update of the variable

# List-Directed I/O (3)

As described, **slashes** (/) terminates the call  
That is why **slash** needs to be **quoted**

Before using it in complicated, important code:

- Read the specification, to avoid “**gotchas**”
- Work out exactly what you want to do with it

# Formatted Input for REALs

**m** in **Fn.m** etc. is an implied **decimal point**

It is used **only** if you don't provide one

The **k** in **En.mEk** is completely ignored

And there are more **historical oddities**

Here is an **extended** set of rules

- Use a **precision** of zero (e.g. **F8.0**)
- Always include a **decimal point** in the number
- Don't use the **P** or **BZ** **descriptors** for **input**
- Don't set **BLANK='zero'** in **OPEN** or **READ**

# The Sordid Details

If you want to know, read the actual standard  
You won't believe me if I tell you!

And don't trust any books on this matter  
They **all** over-simplify it like crazy

In any case, I doubt that any of you care  
Follow the above rules and you don't need to



# Choice of Unit Number

**Preconnected units** are open at program start

Includes at least ones referred to by **UNIT=\***

- **OPEN** on them will **close** the old connection

Can check for an **open unit** using **INQUIRE**

**Fortran 2003** has a way of getting their **numbers**

Has **names** in the **ISO\_FORTRAN\_ENV** module

Critical only for significant, portable programs

# INQUIRE By File (1)

You can check if a file **exists** or is **open**

LOGICAL :: here

INQUIRE (FILE='name', EXIST=here)

INQUIRE (FILE='name', OPENED=here)

- These answers may not mean what you expect  
E.g. a new, output file may be **open** but not **exist**
- Name matching may be **textual** or by **identity**  
Watch out when using **ln** or **ln -s**

# INQUIRE By File (2)

Can query **SIZE**, **READ**, **READWRITE**, **WRITE**

Don't bet on it – not all compilers support them sanely

Some others, too, but not under **Unix-like** systems

Most other queries are handled like **inquire by unit**

Subject to matching the **file name** correctly

If **not connected** always return **UNKNOWN**

Not exactly the most useful behaviour!

However, at least they **DO** say **UNKNOWN**

And don't simply return plausible nonsense

# INQUIRE By Unit (1)

**Inquire by unit** most usefully does two things:

Checks if the **unit** is currently **connected**

Returns the **record length** of an open file

```
LOGICAL :: connected
```

```
INQUIRE (UNIT=number, OPENED=connected)
```

```
INTEGER :: length
```

```
INQUIRE (UNIT=number, RECL=length)
```

You can ask about both together, of course

# INQUIRE By Unit (2)

There are other potentially useful specifiers  
Not all of them make much sense under **POSIX**

You can get all of the **specifiers** used for **OPEN**  
Could be useful when writing generic libraries

**SIZE** gives the size of the file, probably in bytes  
This is only in **Fortran 2003**, and unreliable  
Again, nothing to do with Fortran, as such

See the references for details on them

# Unformatted I/O

Using **pipes** or **sockets** is unreliable

The reasons are complicated and historical

So is **unformatted** I/O of **derived types**

The same applies in **C++**, for very similar reasons

- Ask for advice if you need to do these

# Namelist

**Namelist** is a historical oddity, new in **Fortran 90**  
This sounds impossible, but I assure you is true

- Not recommended, but not deprecated, either

# STREAM Files

Fortran 2003 has introduced **STREAM** files  
These are for interchange with **C**-like files  
They provide all **portable** features of **C**

- They allow **positioning**, like **C** text files  
I advise avoiding the **POS=** specifier  
It's full of **gotchas** in both **C** and Fortran



# I/O of Derived Types

The **DT** descriptor has been mentioned

- Unfortunately, it's often **not implemented**

You can do almost anything you need to  
But this course cannot cover everything

# Asynchronous I/O

Mainframes proved that it is the right approach  
Fortran 2003 introduced it

- For complicated reasons, you should avoid it
- This has **nothing** to do with Fortran  
Don't use **POSIX** asynchronous I/O, either  
And probably not **Microsoft's** . . .

# BACKSPACE

Don't go there

It was provided for magnetic tape file support  
In those days, could often read backwards, too

It's almost always a performance disaster, at best  
And it very often doesn't actually work reliably

- Again, that is NOT specific to Fortran  
It applies to using seek in C/C++, too  
Never reposition on sequential files  
Rewinding to the beginning is usually OK

# Oddities of Connection

- Try to avoid these, as they are confusing  
You will see them in some of the references

**Files** can be **connected** but not **exist**

Ones newly created by **OPEN** may be like that

**Units** can be **connected** when the program starts

Ask me if you want to know why and how

**OPEN** can be used on an existing **connection**

It modifies the connection properties

# Other Topics

There are a **lot** more optional features  
You must read Fortran's specifications for them

**Fortran 2003** adds many slightly useful features  
Most compilers don't support many of them yet  
The above has described the most useful ones

And a **few** features should be avoided entirely

For more on this, look at the **OldFortran** course

# Last Reminder

Be careful when using Fortran I/O features  
They don't always do what you expect

It is much cleaner than C/POSIX, but . . .

Fortran's **model** is very unlike C/POSIX's  
Fortran's **terminology** can be very odd

The underlying C/POSIX can show through  
In addition to Fortran's own oddities